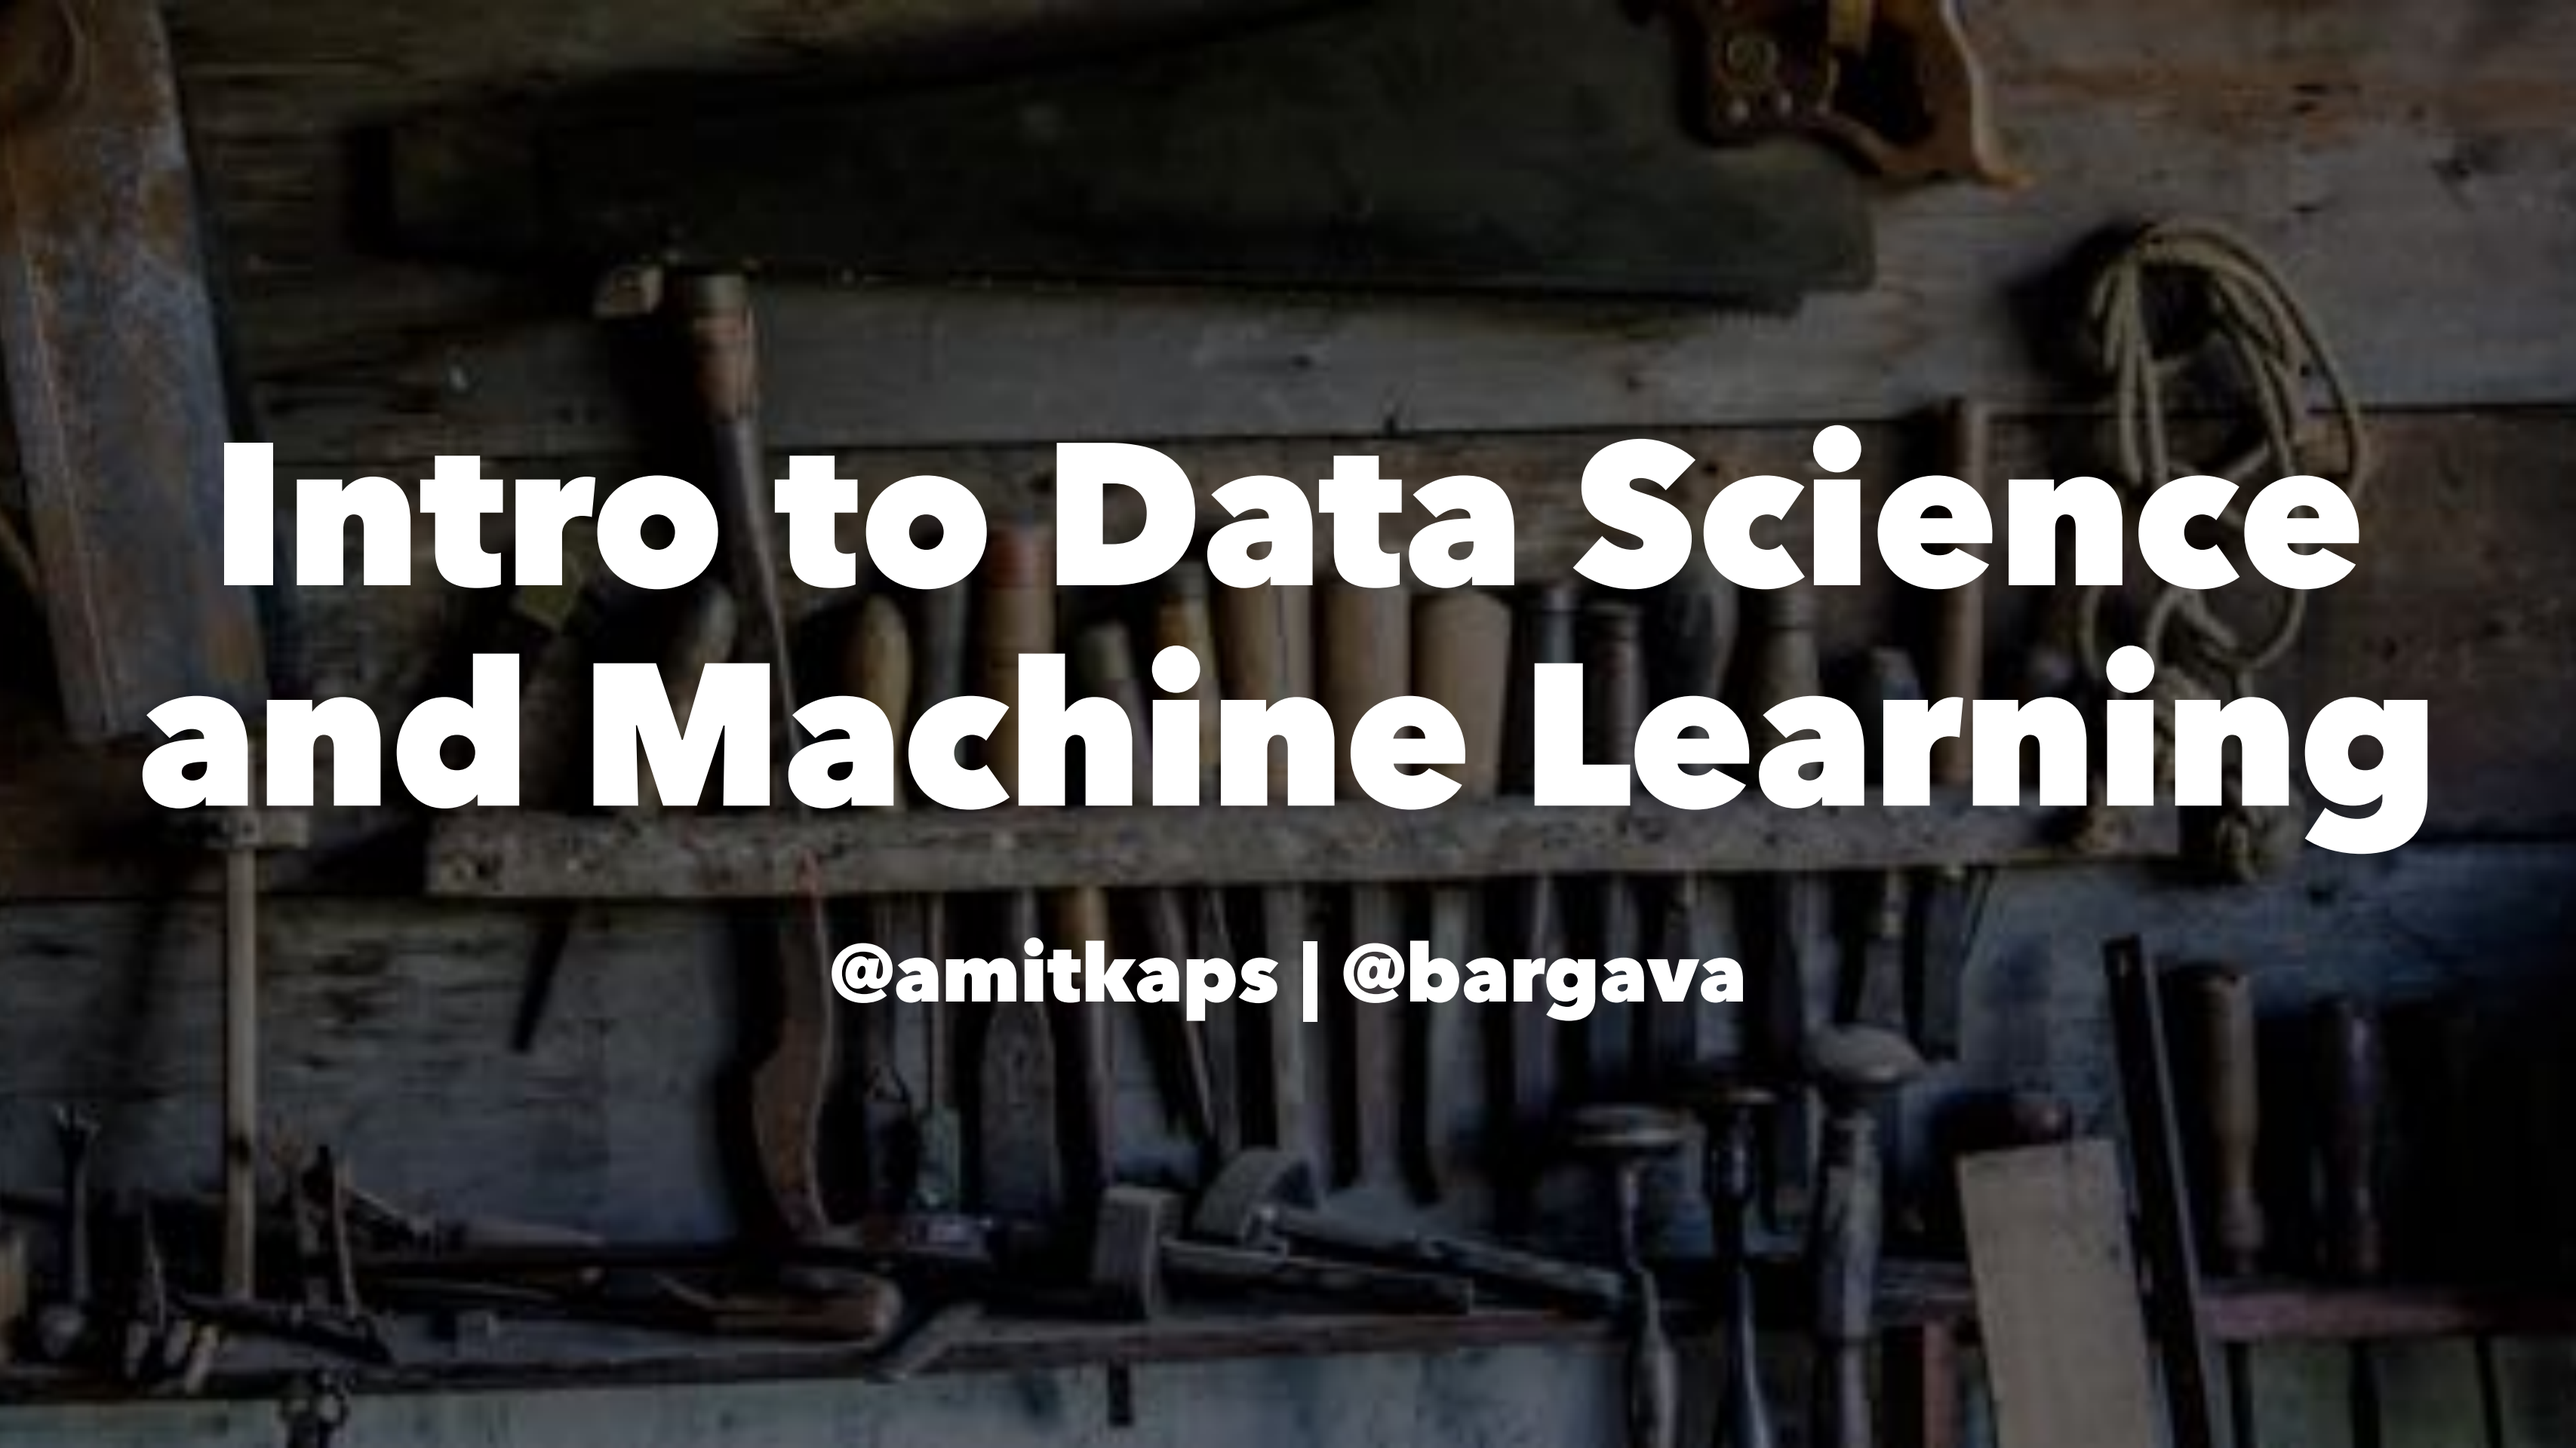

Intro to Data Science and Machine Learning

@amitkaps | @bargava

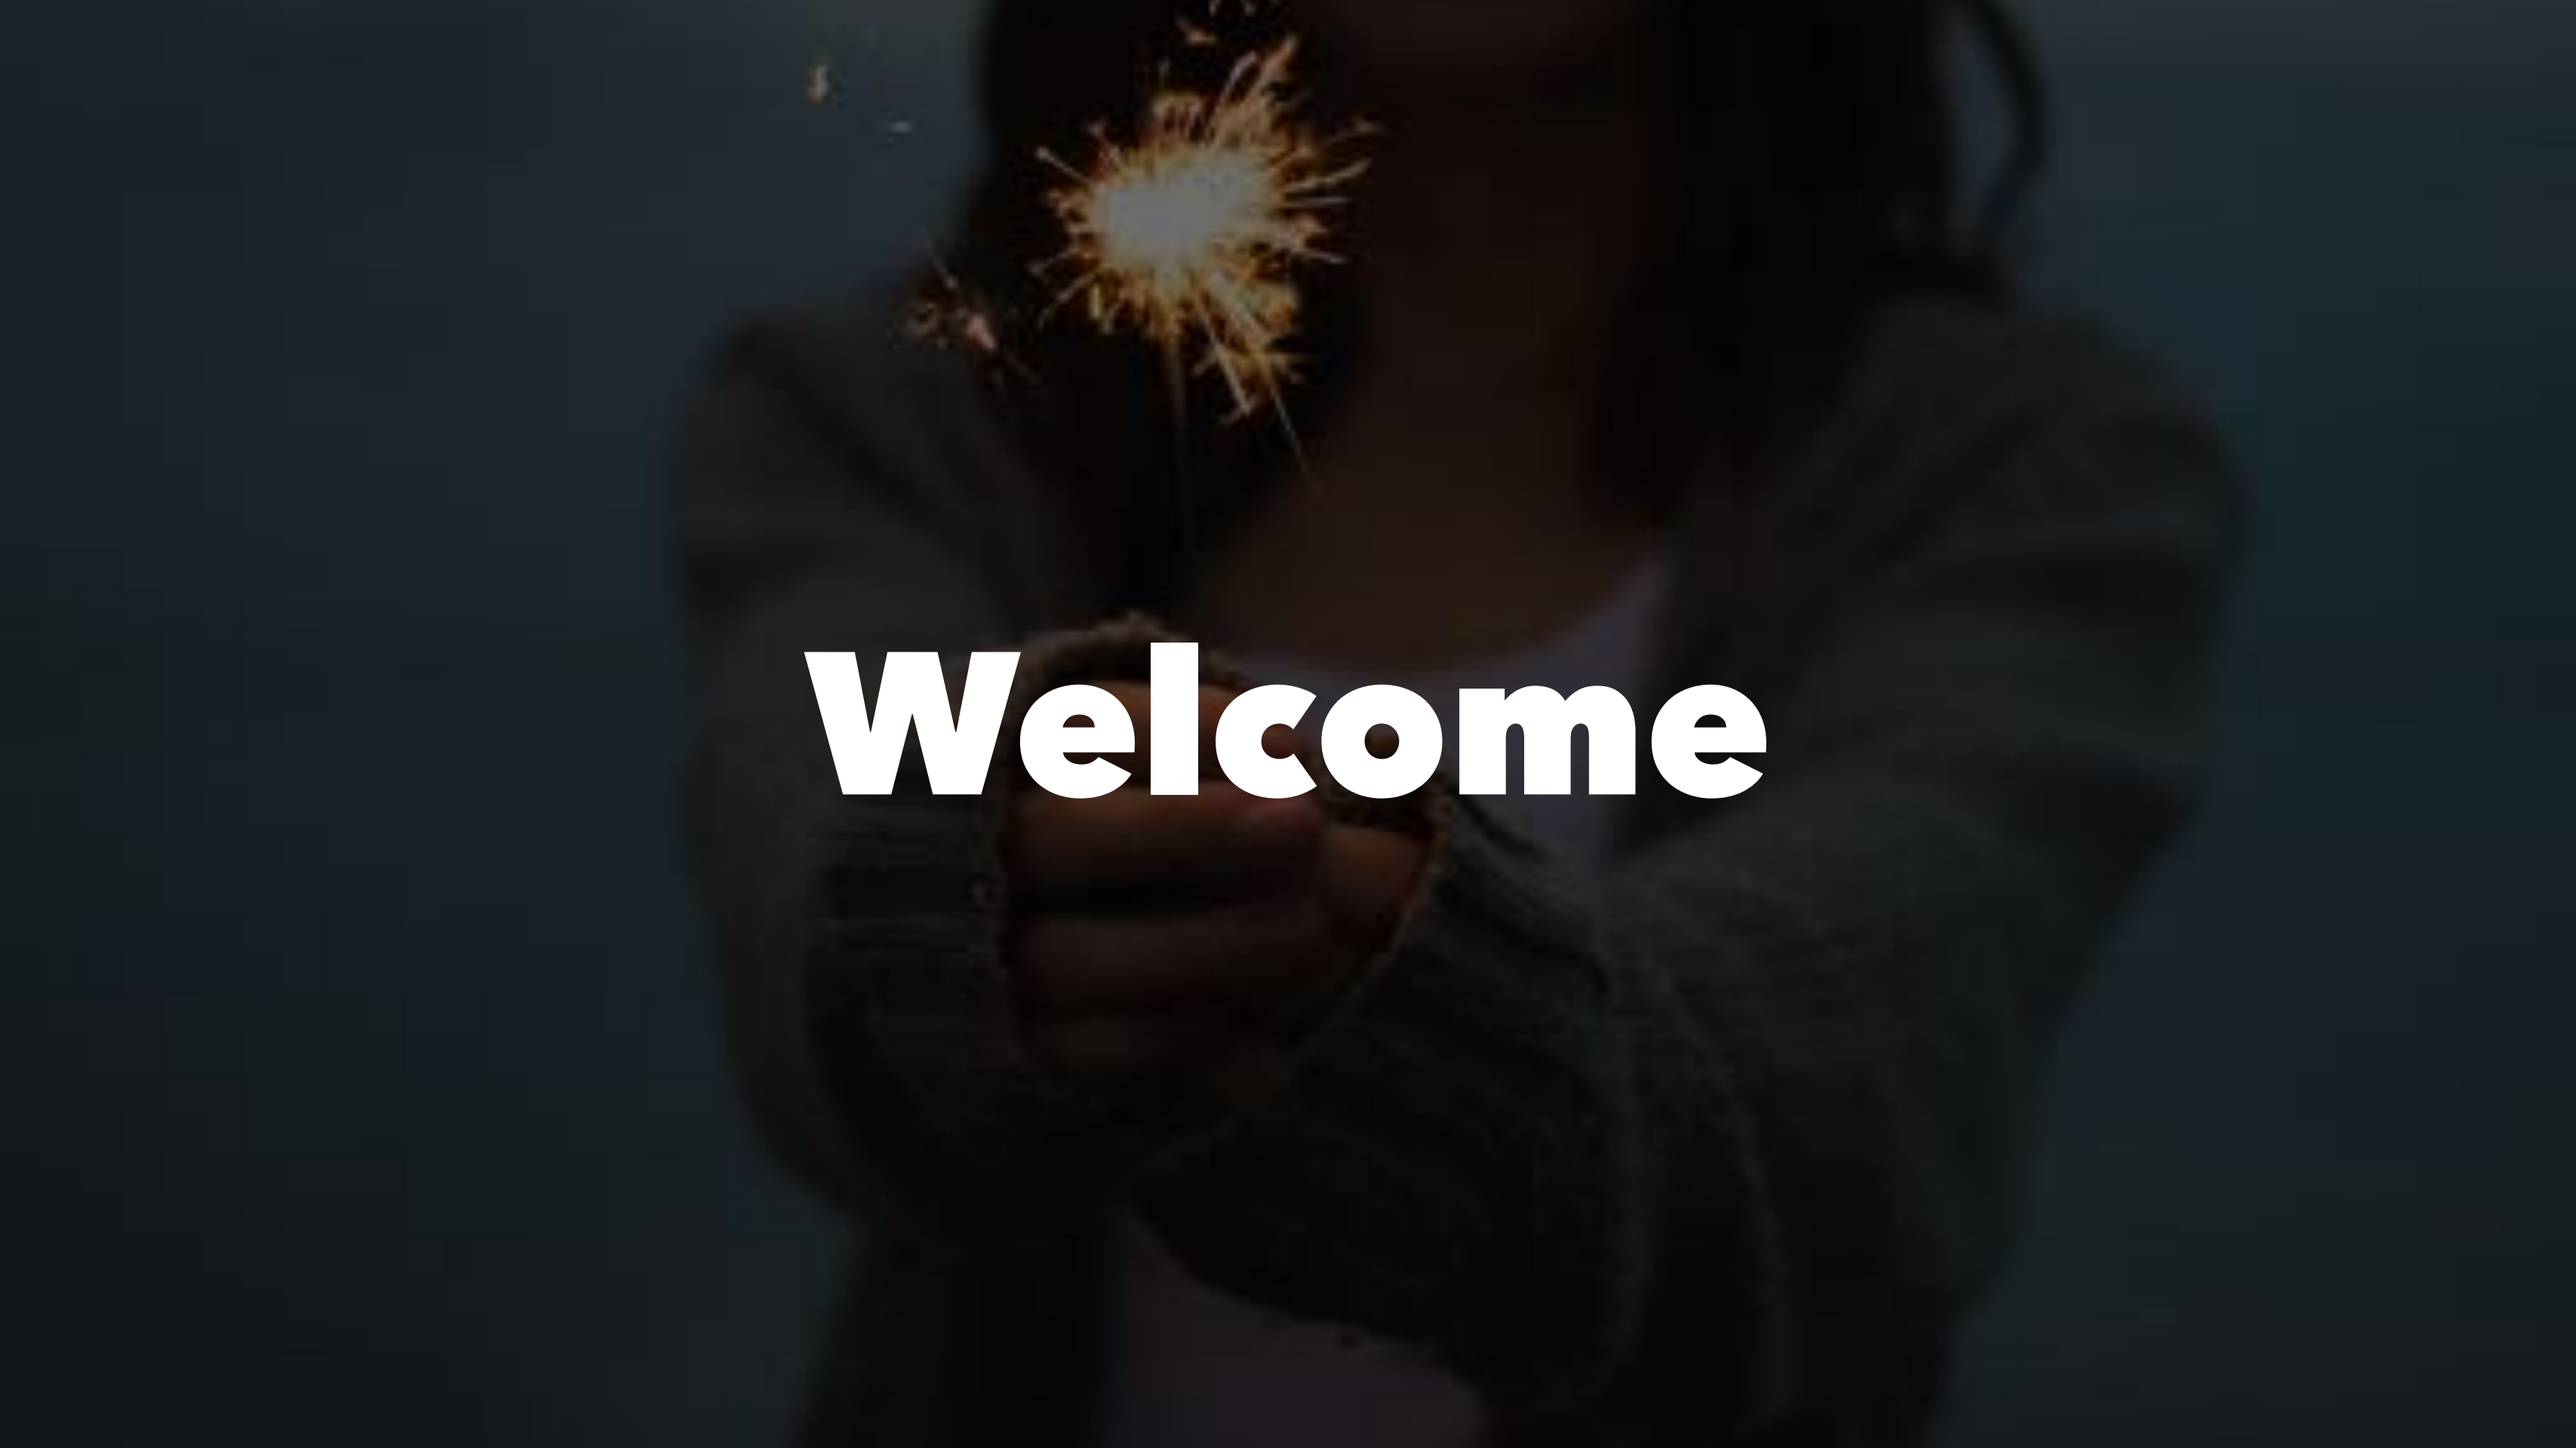
Welcome

Facilitators


Amit
@amitkaps


Bargava
@bargava

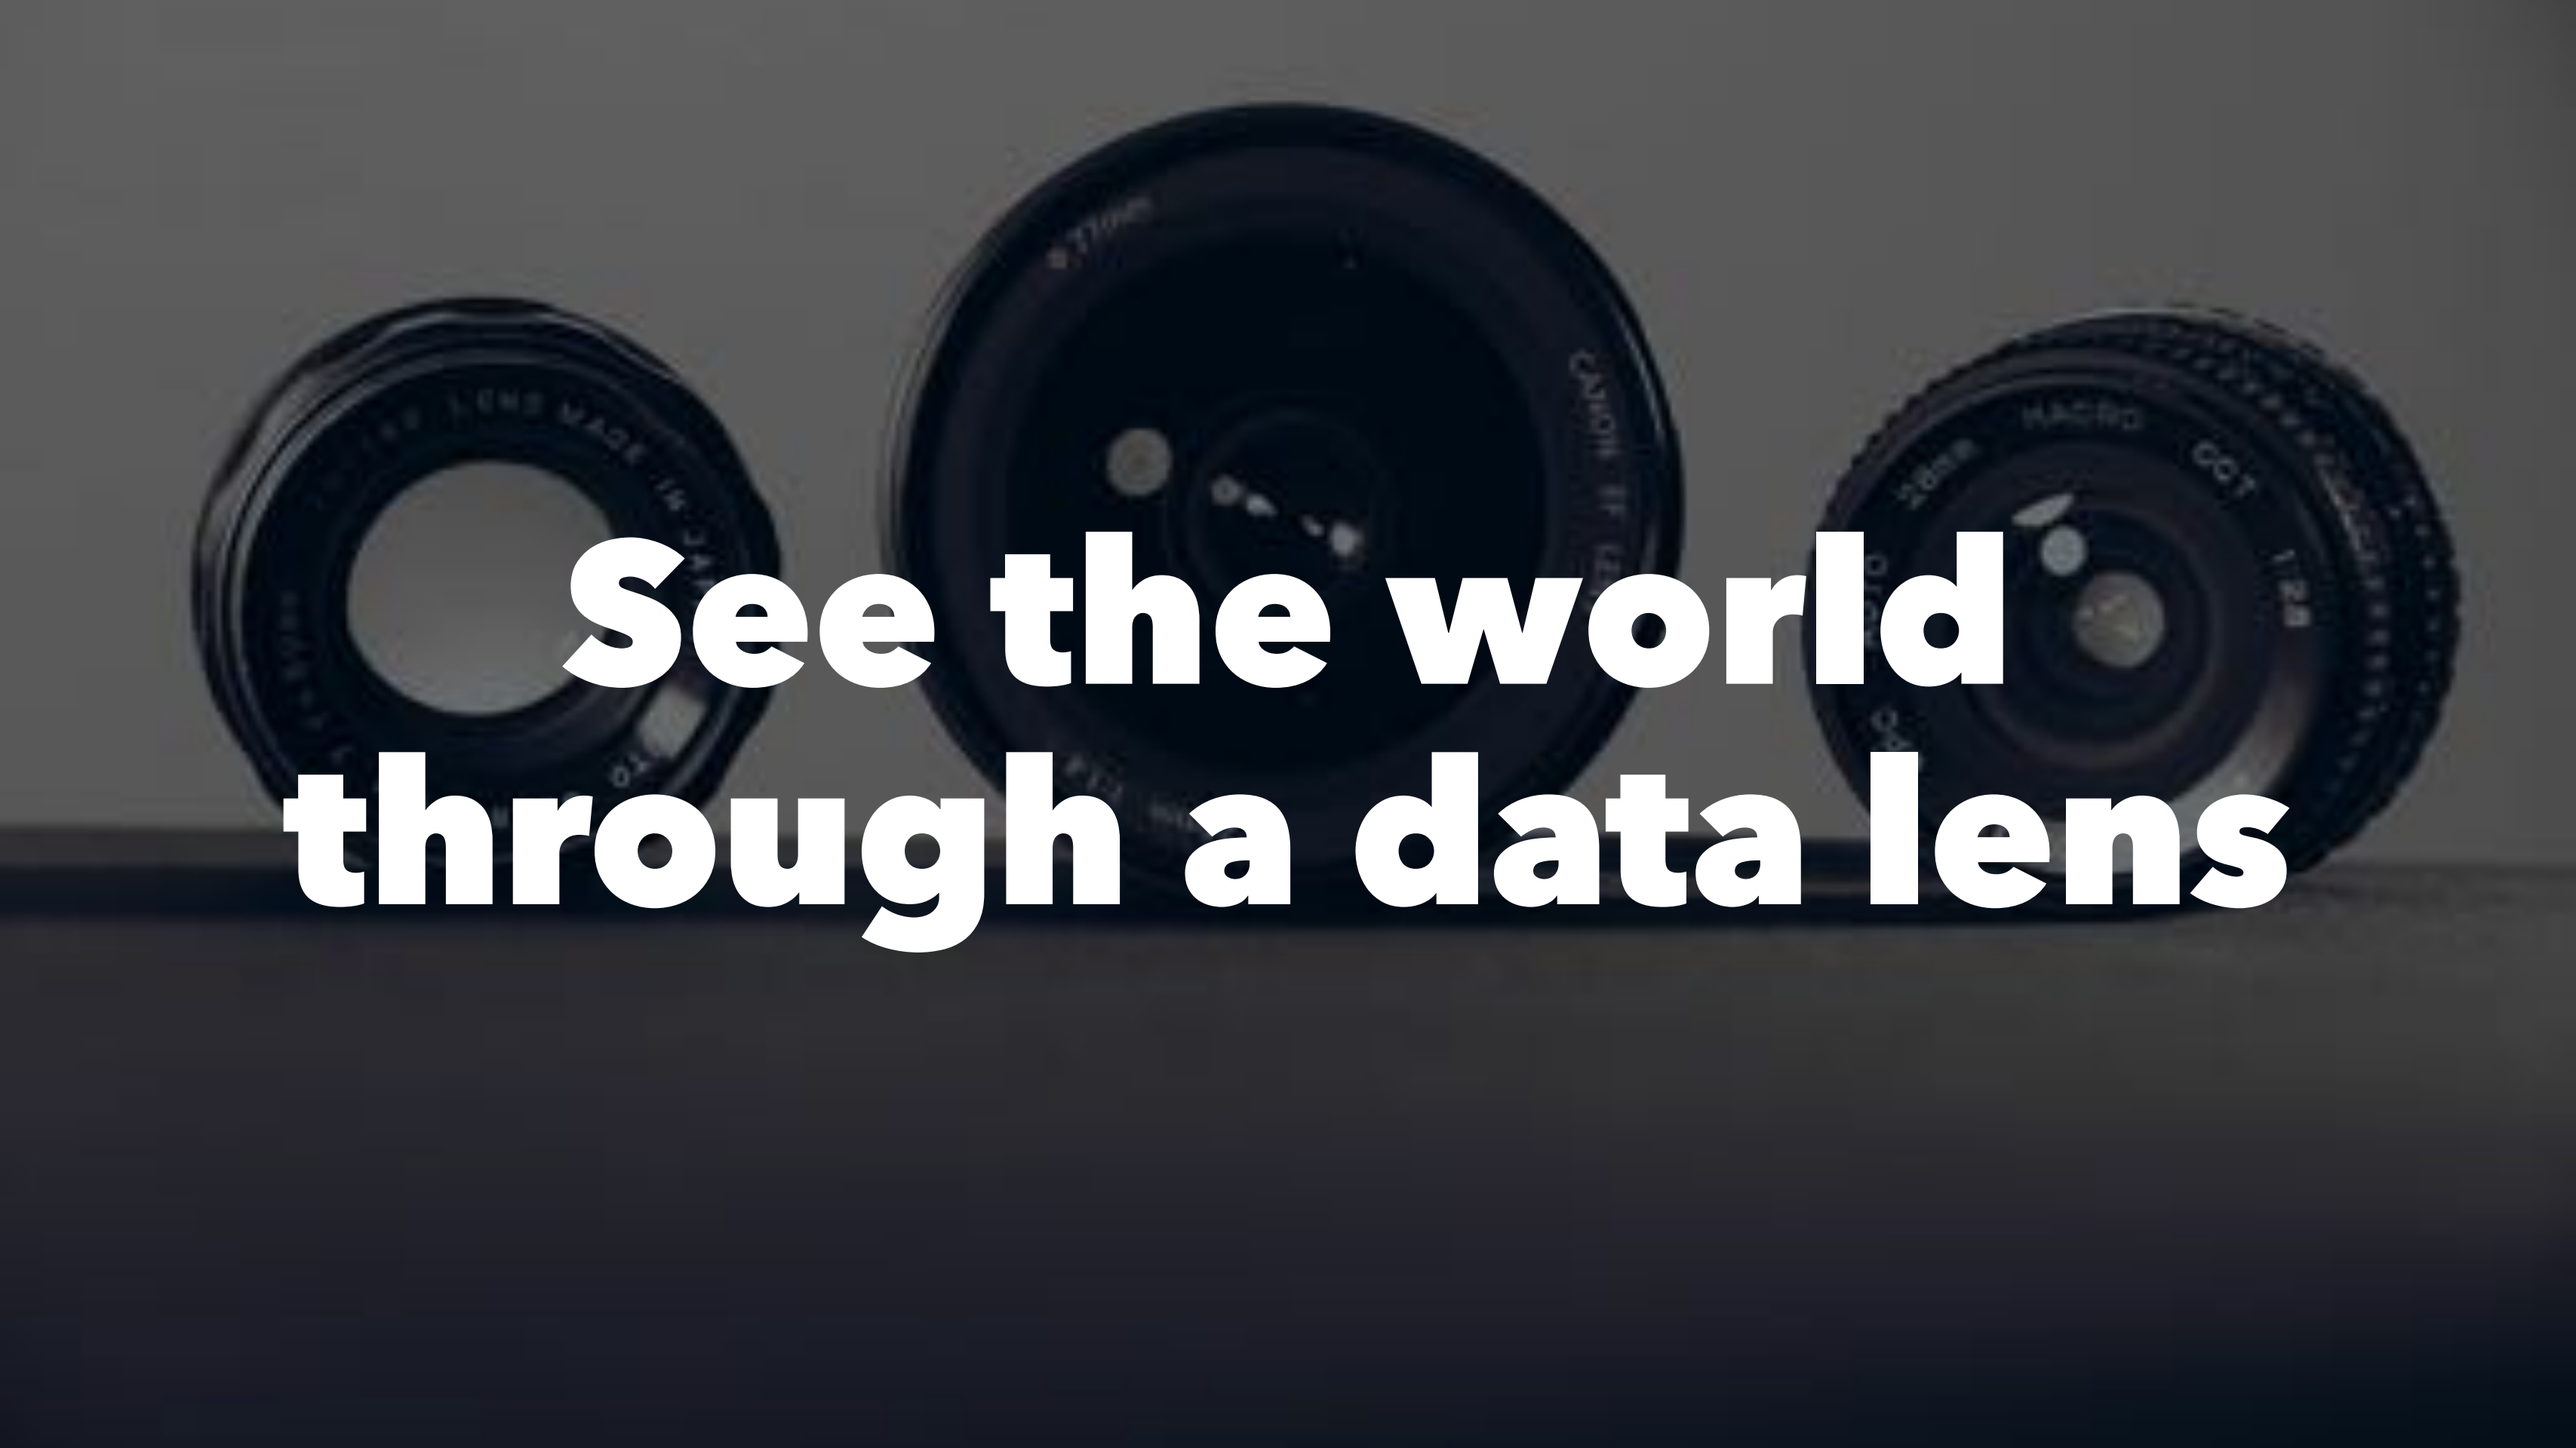
See the world
through a data lens

"Data is just a clue to the end truth"

-- Josh Smith

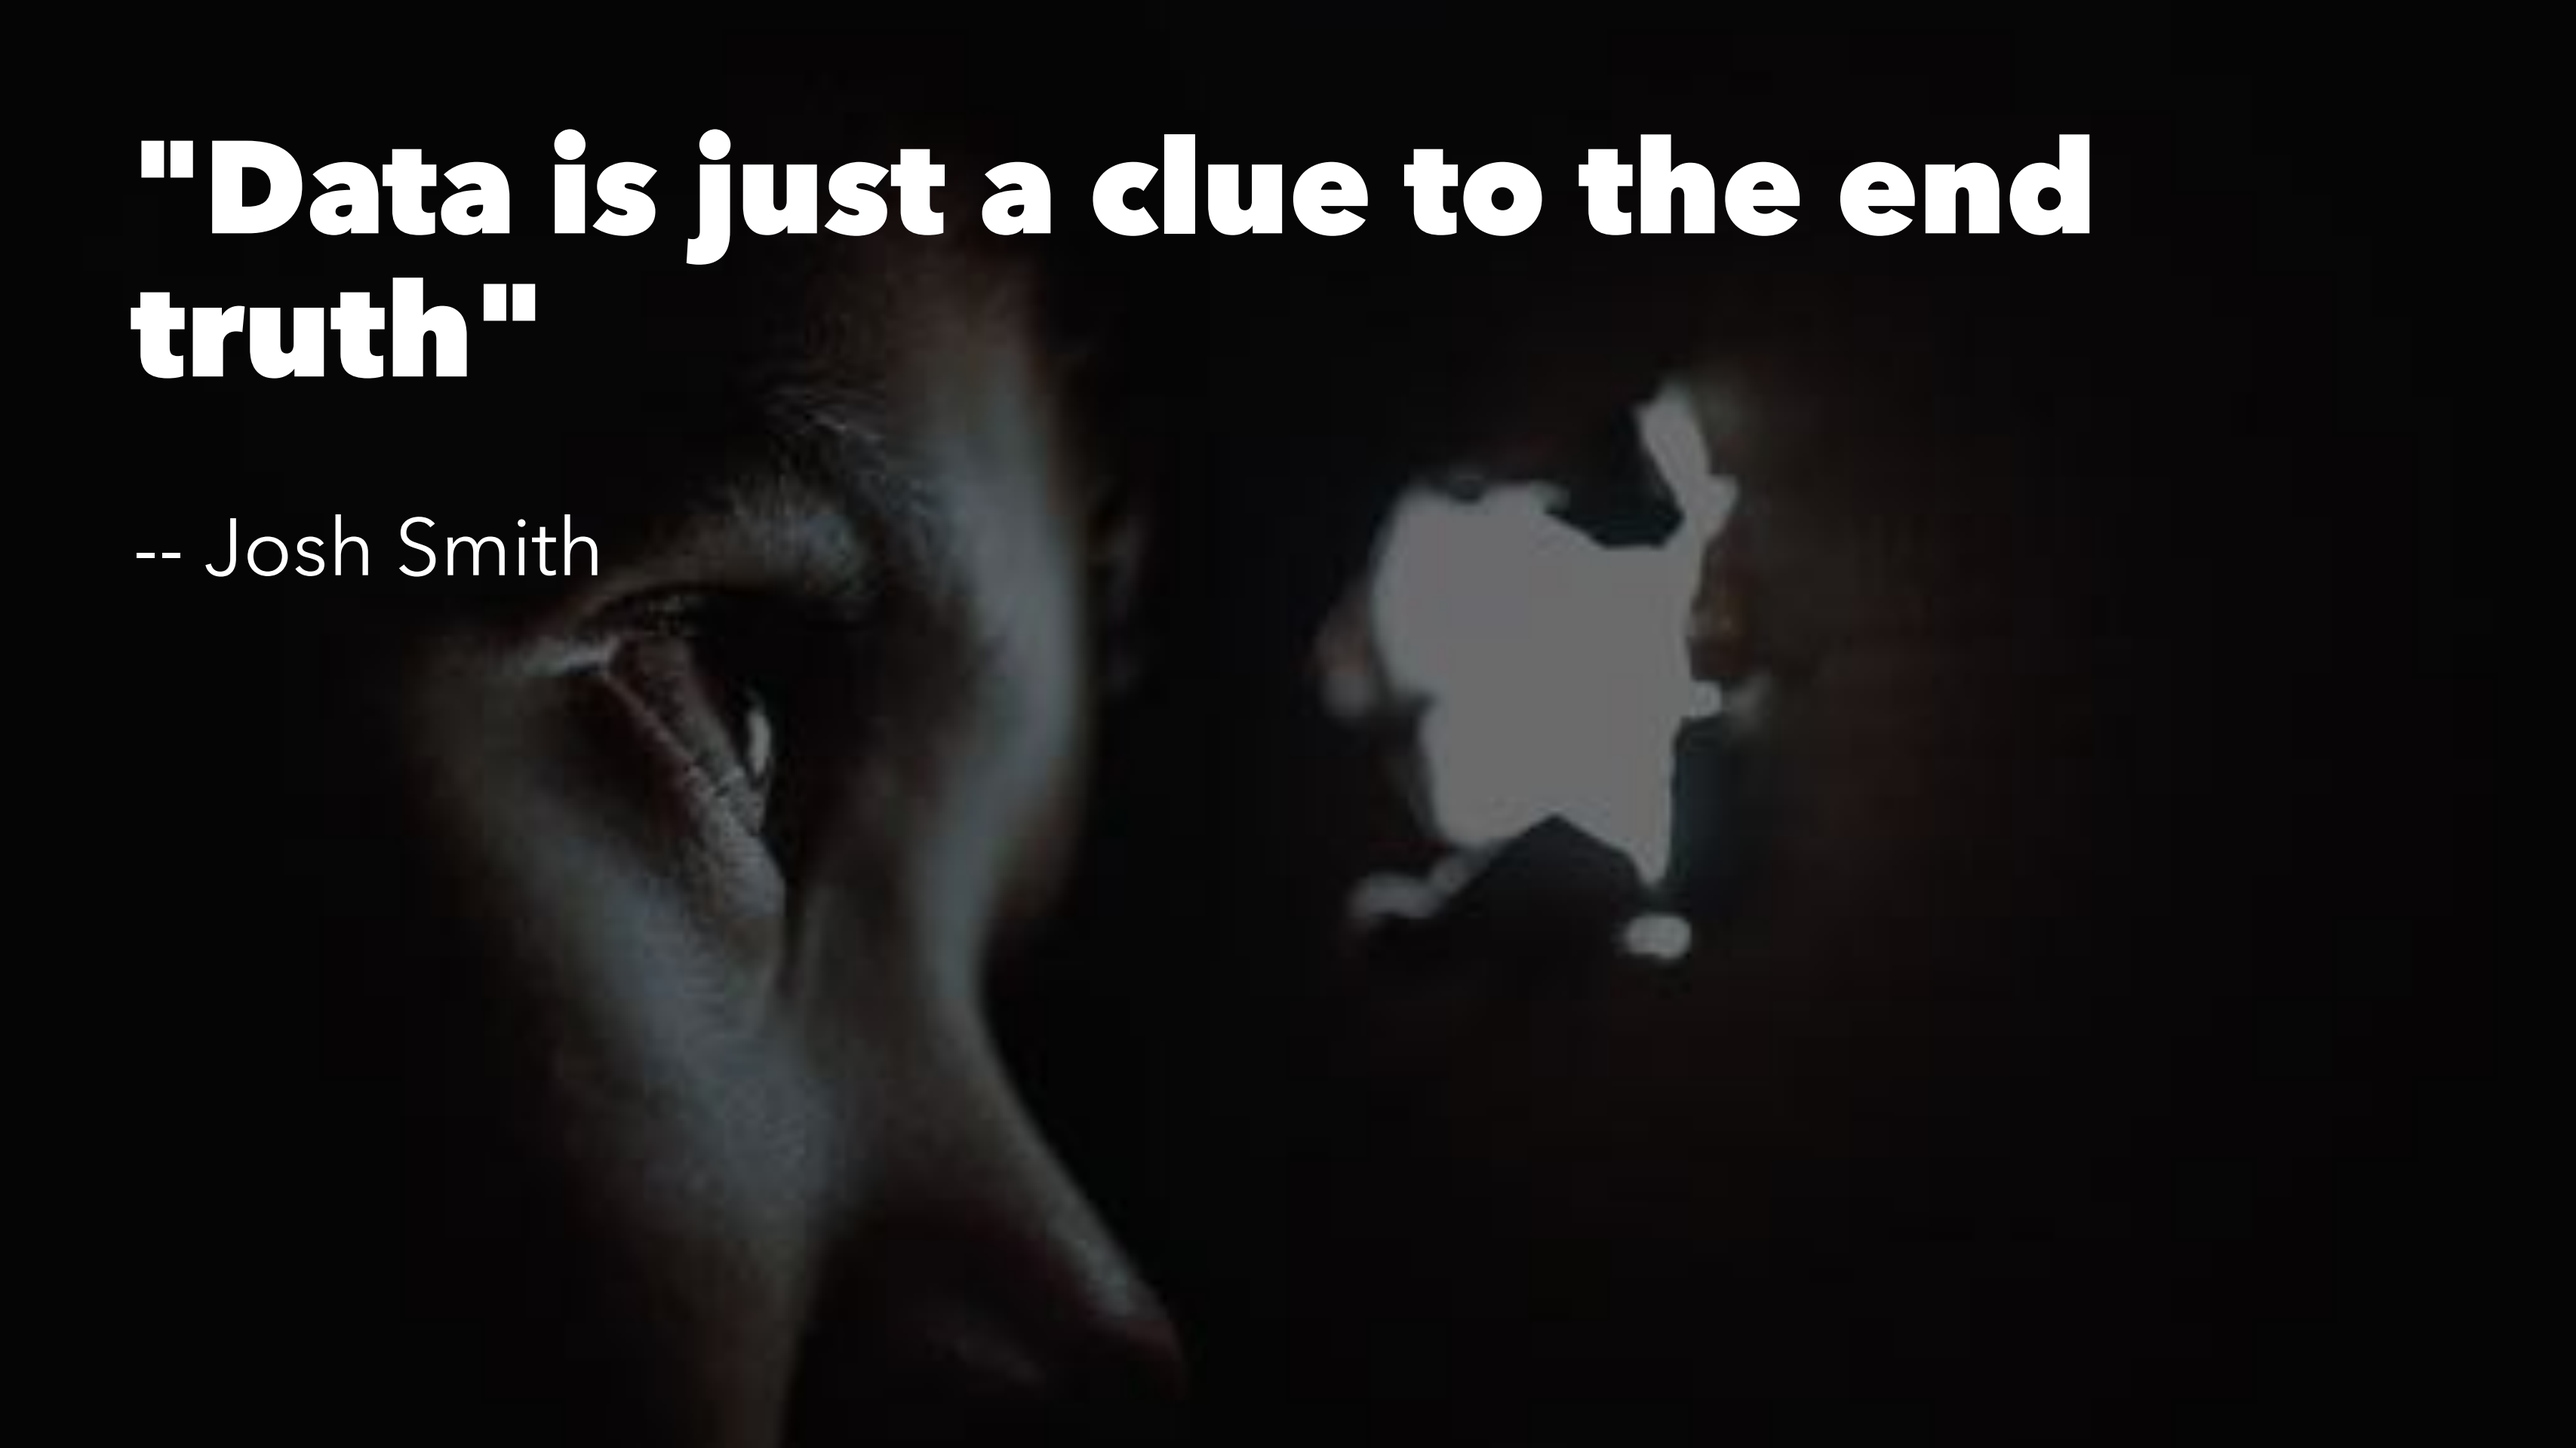


Data Driven Decisions

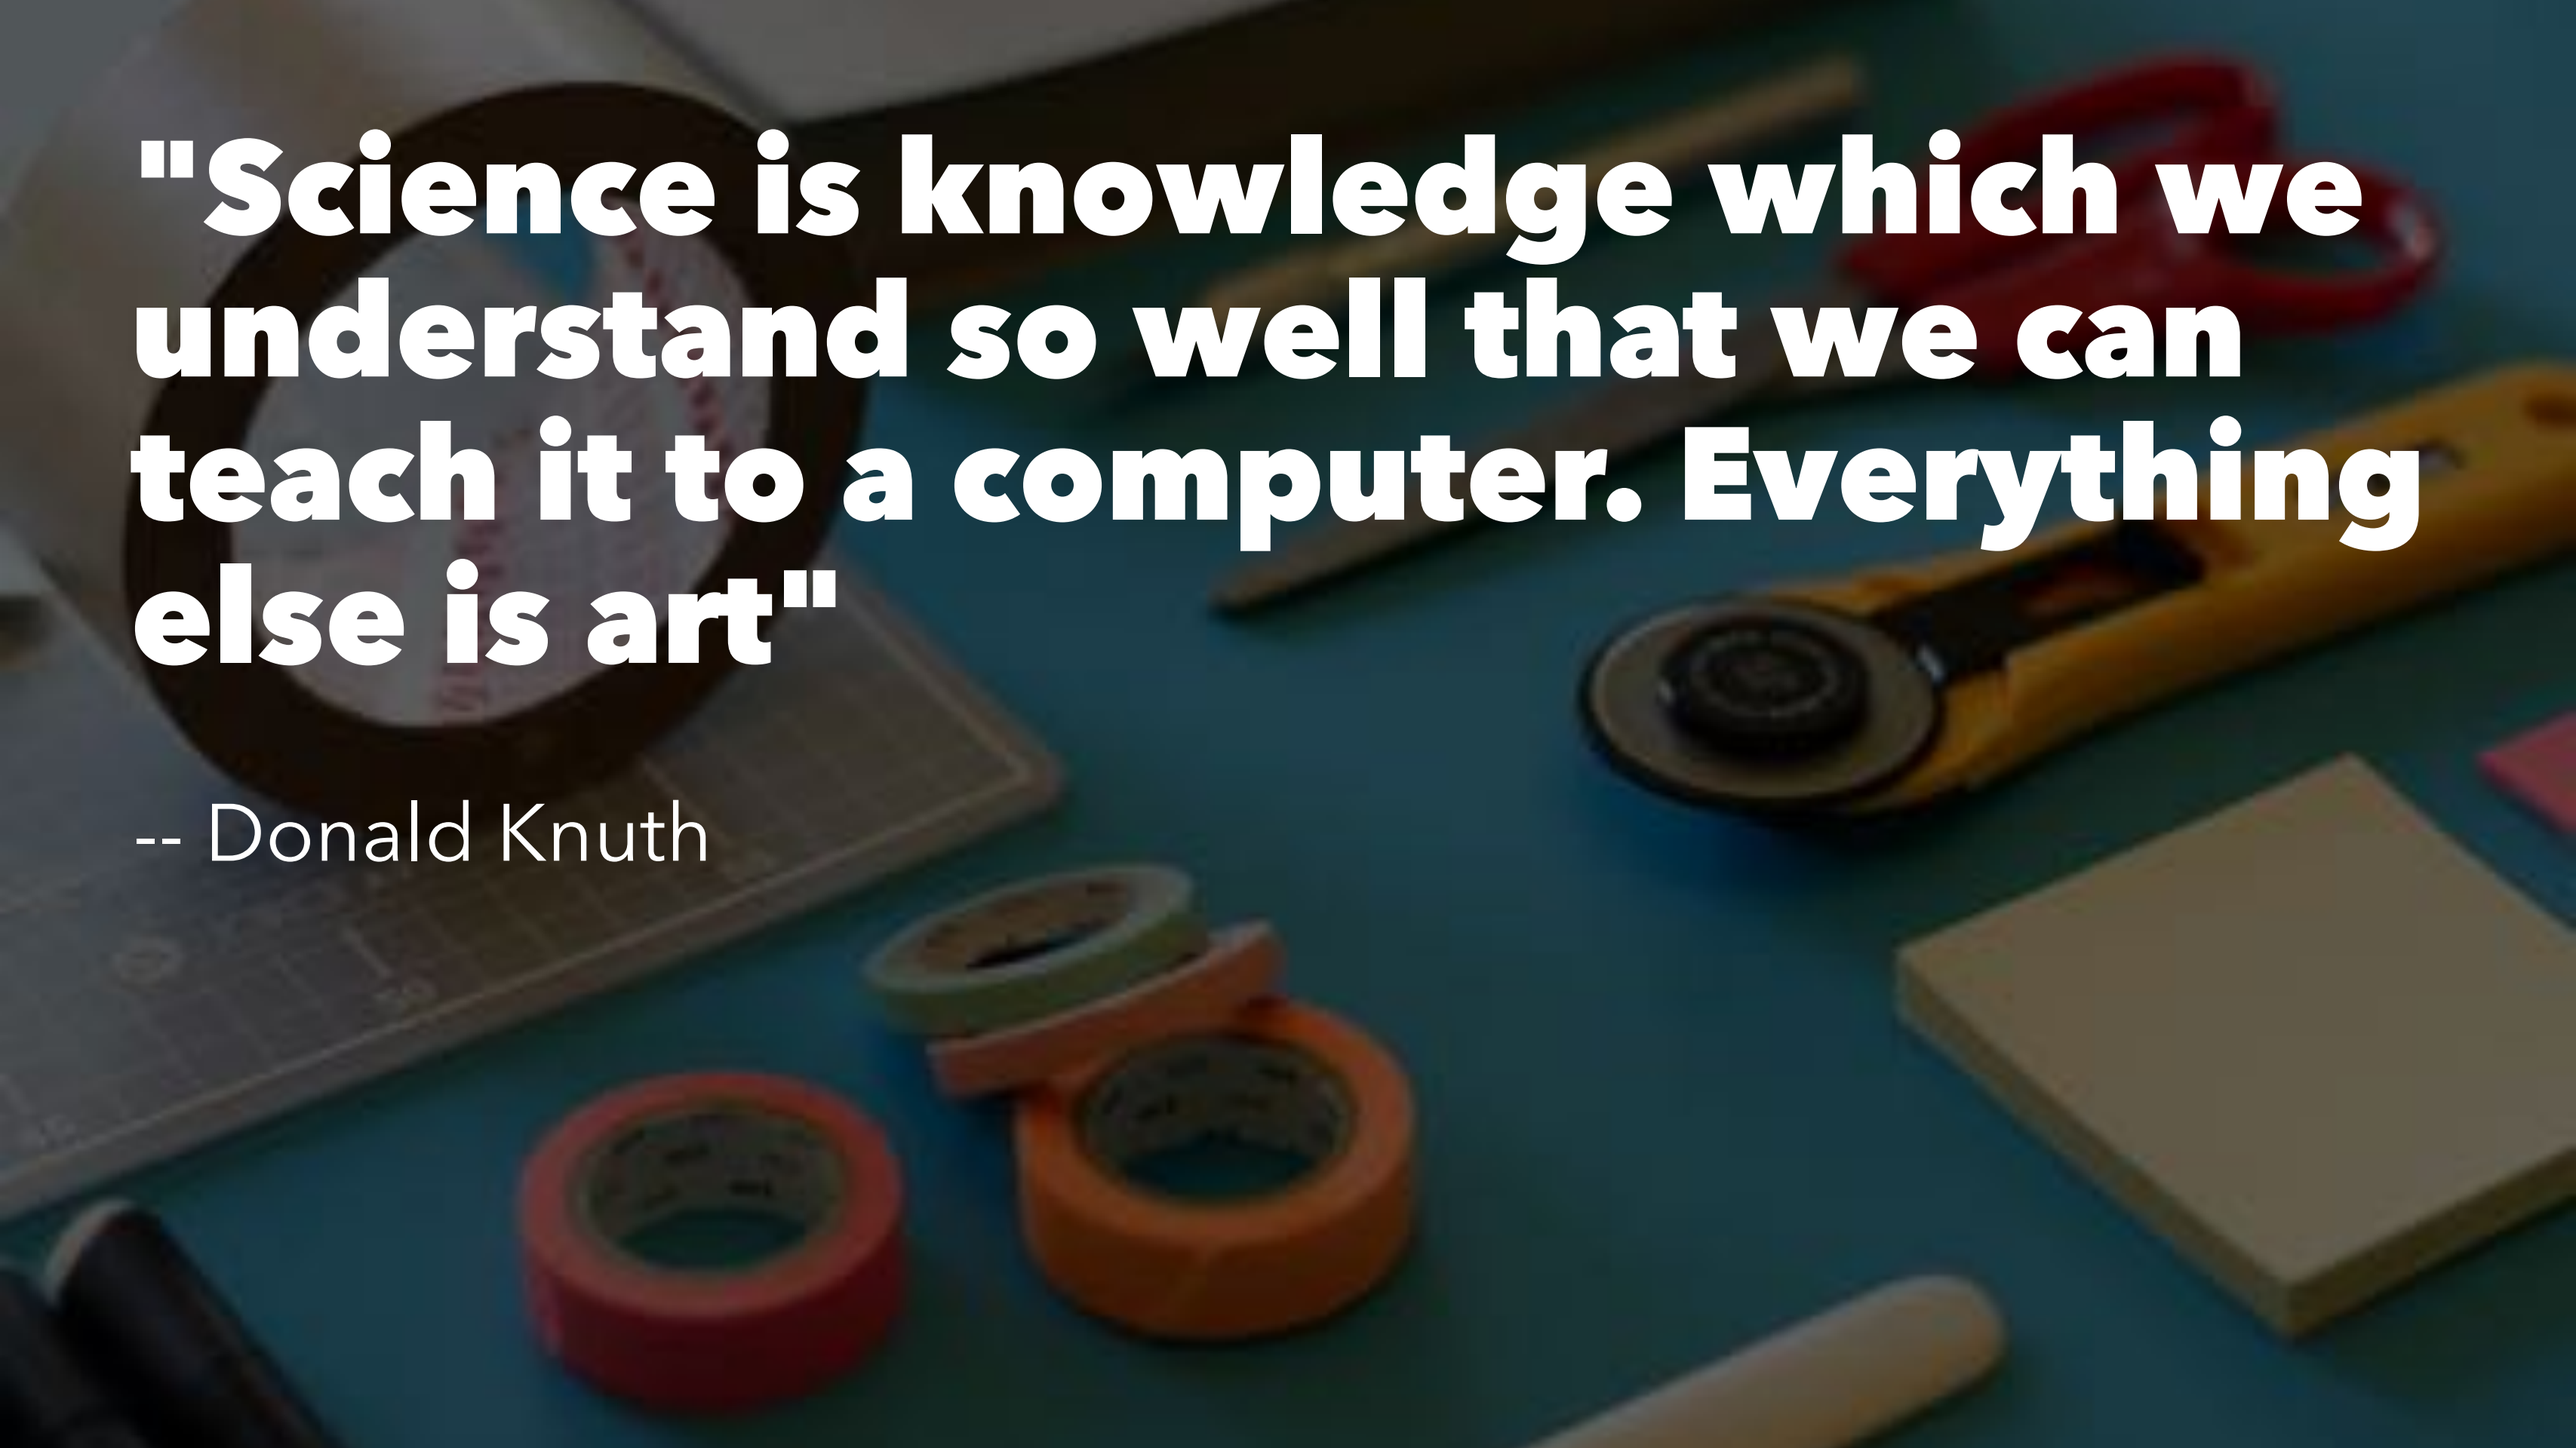
"Science is knowledge which we understand so well that we can teach it to a computer. Everything else is art"
-- Donald Knuth


Data Science is an Art

Hypothesis Driven Approach

Frame

"An approximate answer to the right problem is worth a good deal"

Acquire

"80% perspiration, 10% great idea, 10% great output"


Refine
"All data is messy."

Explore

"I don't know, what I don't know."

Model

"All models are wrong, but some are useful"


Insight
"The goal is to turn data into insight"

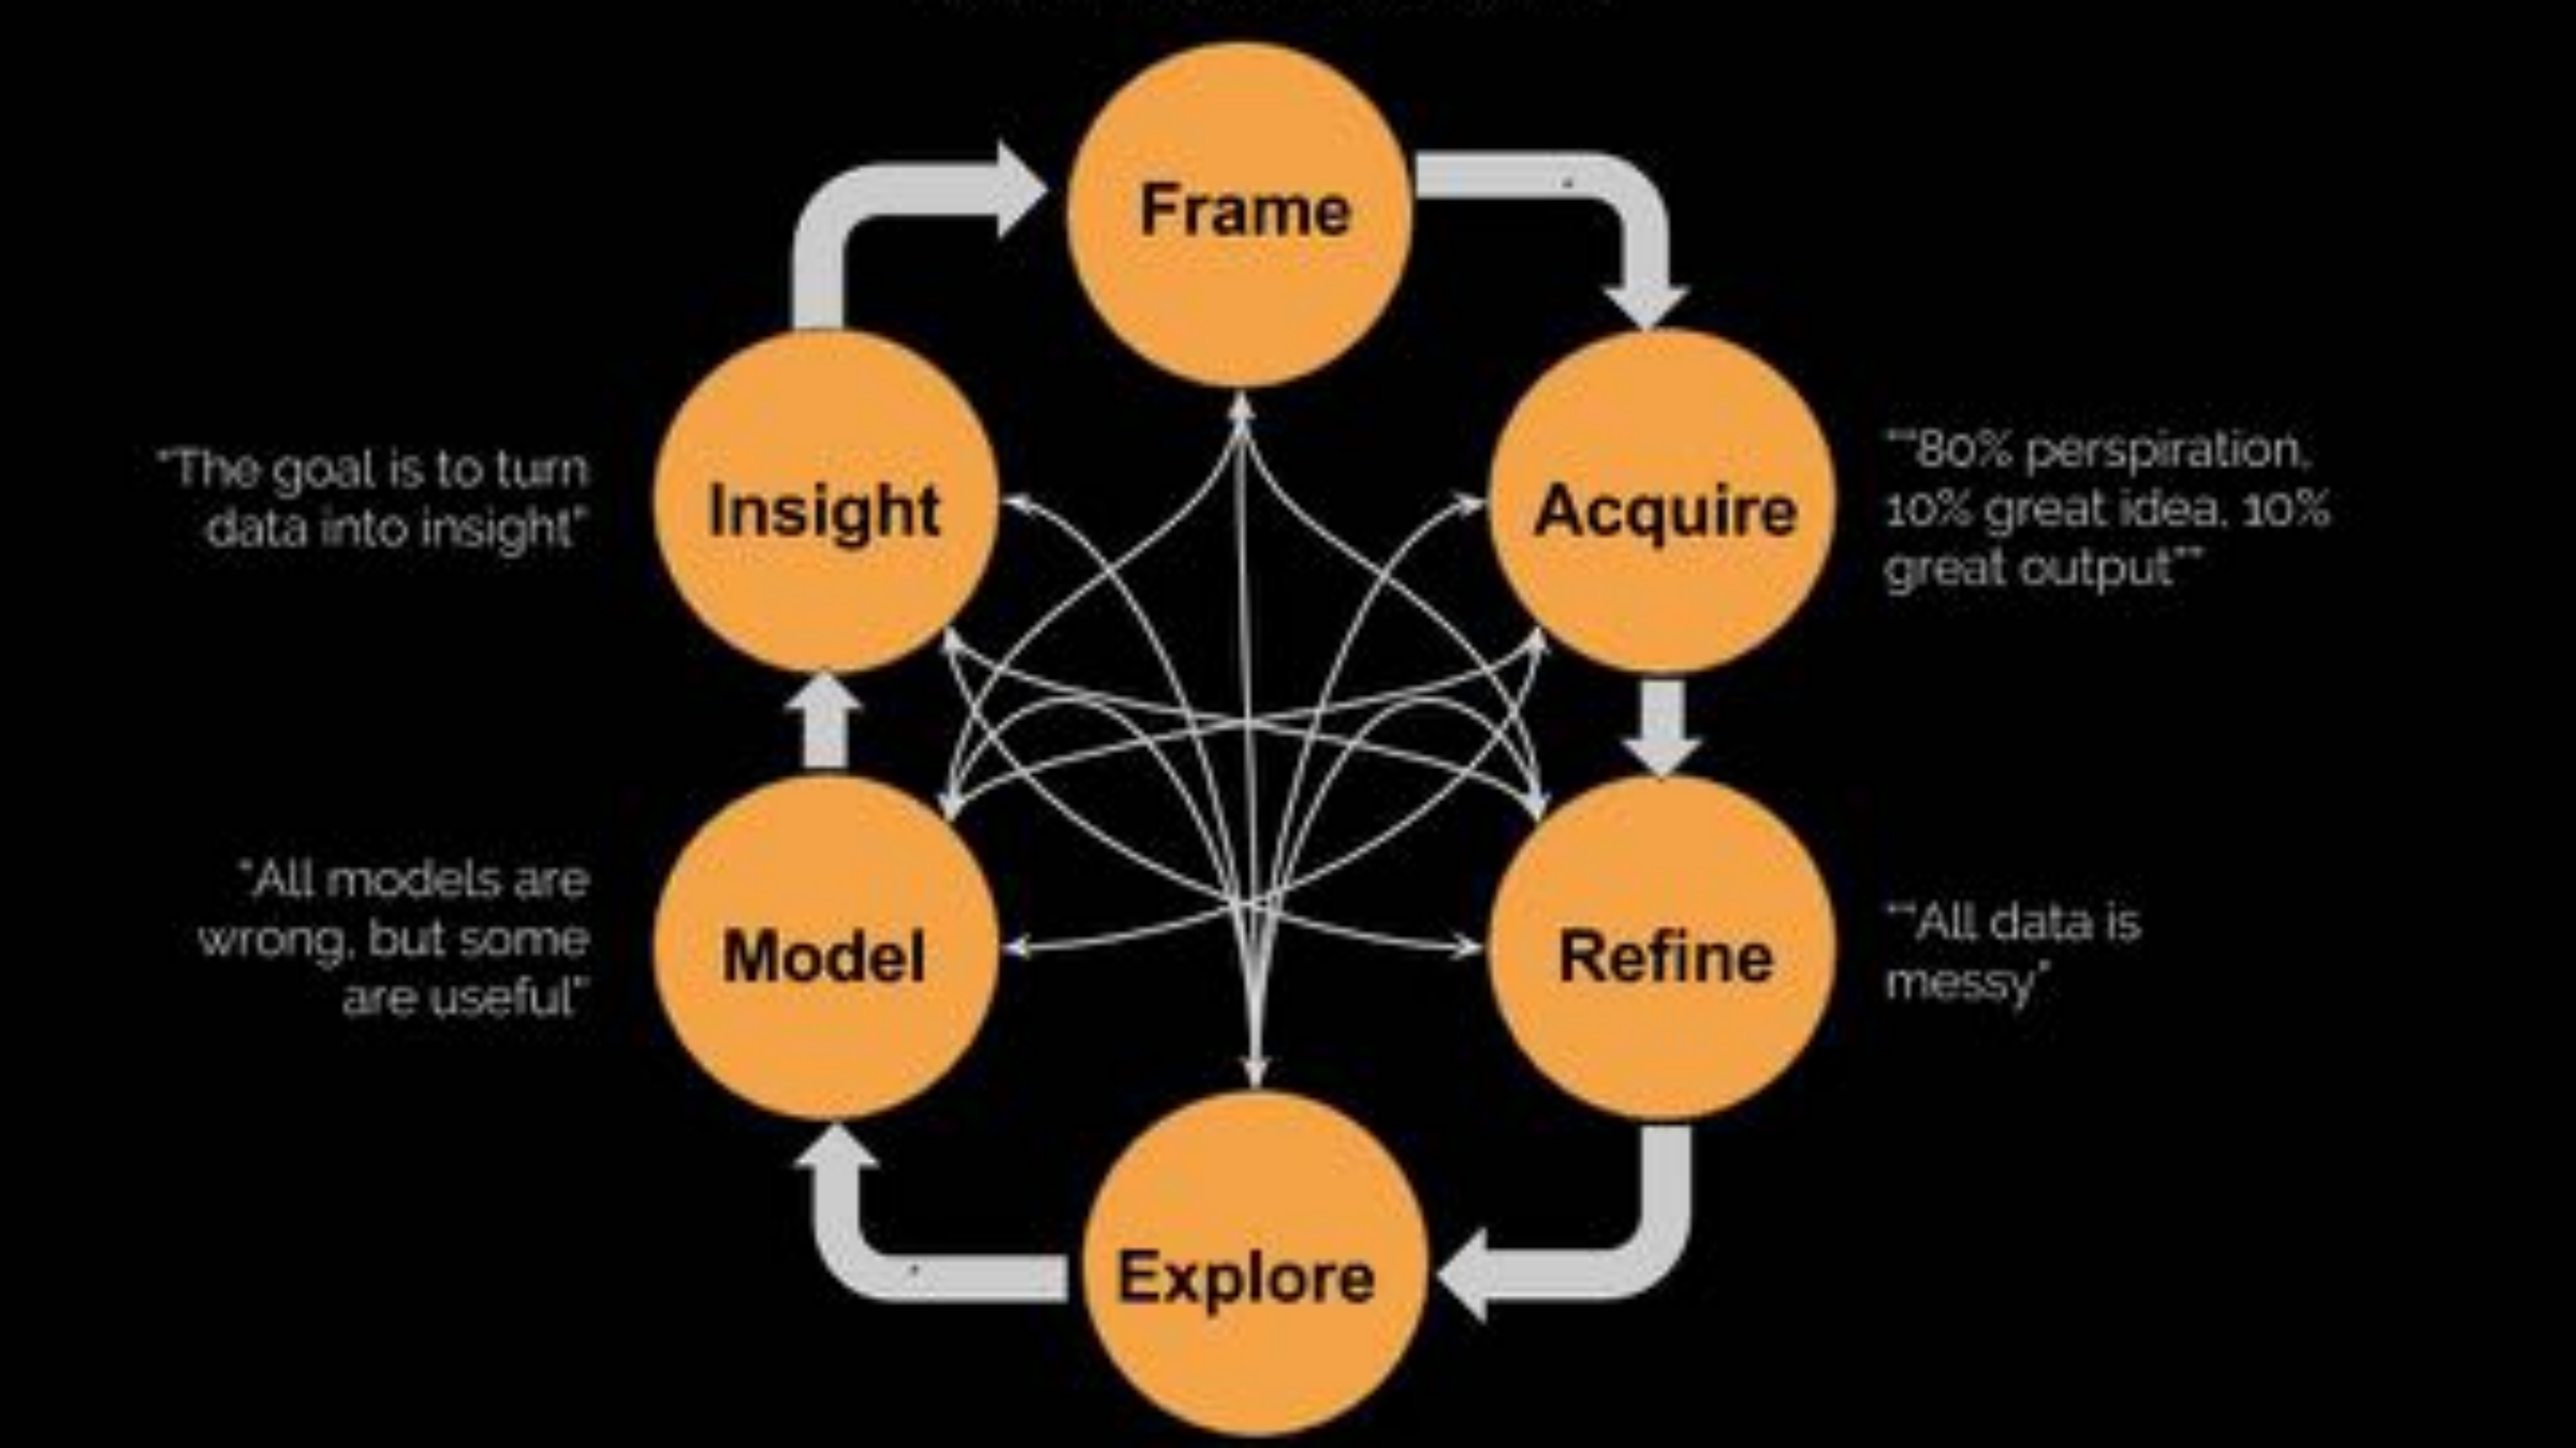
Frame
Acquire
Refine
Explore
Model
Insight
"The goal is to turn data into insight"
"80% perspiration, 10% great idea, 10% great output"
"All models are wrong, but some are useful"
"All data is messy"

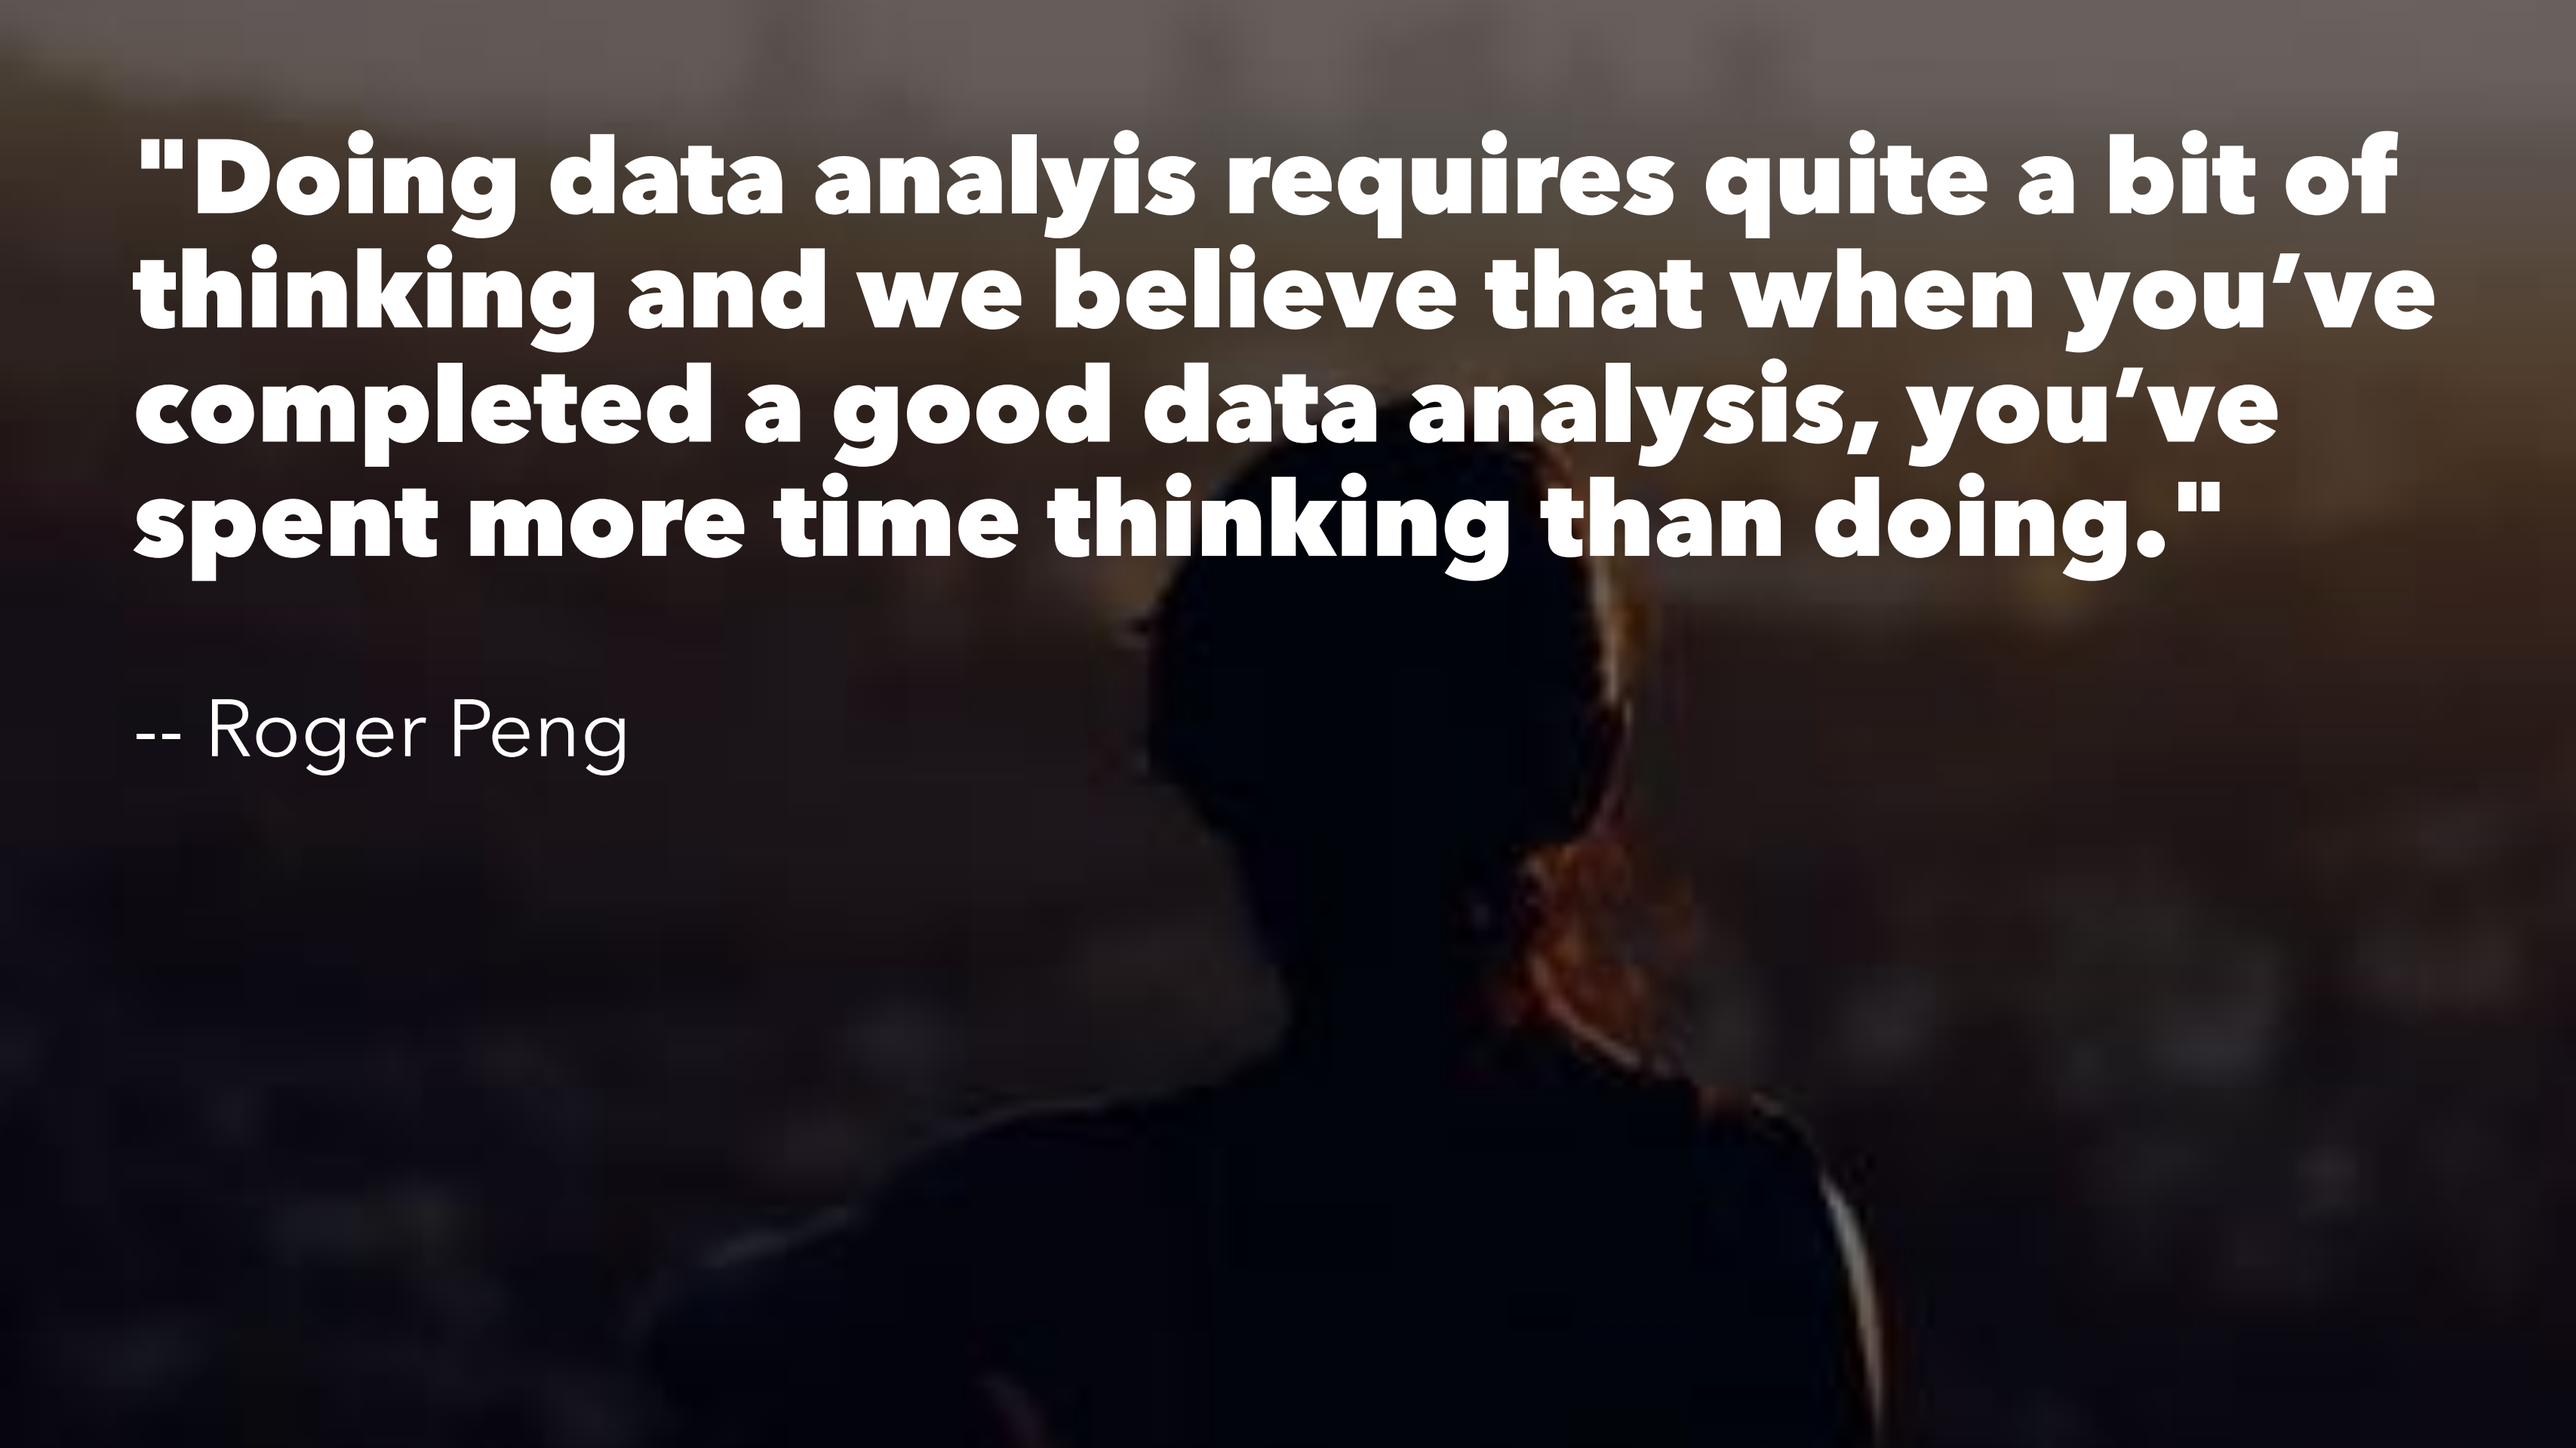
"Doing data analyis requires quite a bit of thinking and we believe that when you've completed a good data analysis, you've spent more time thinking than doing."
-- Roger Peng

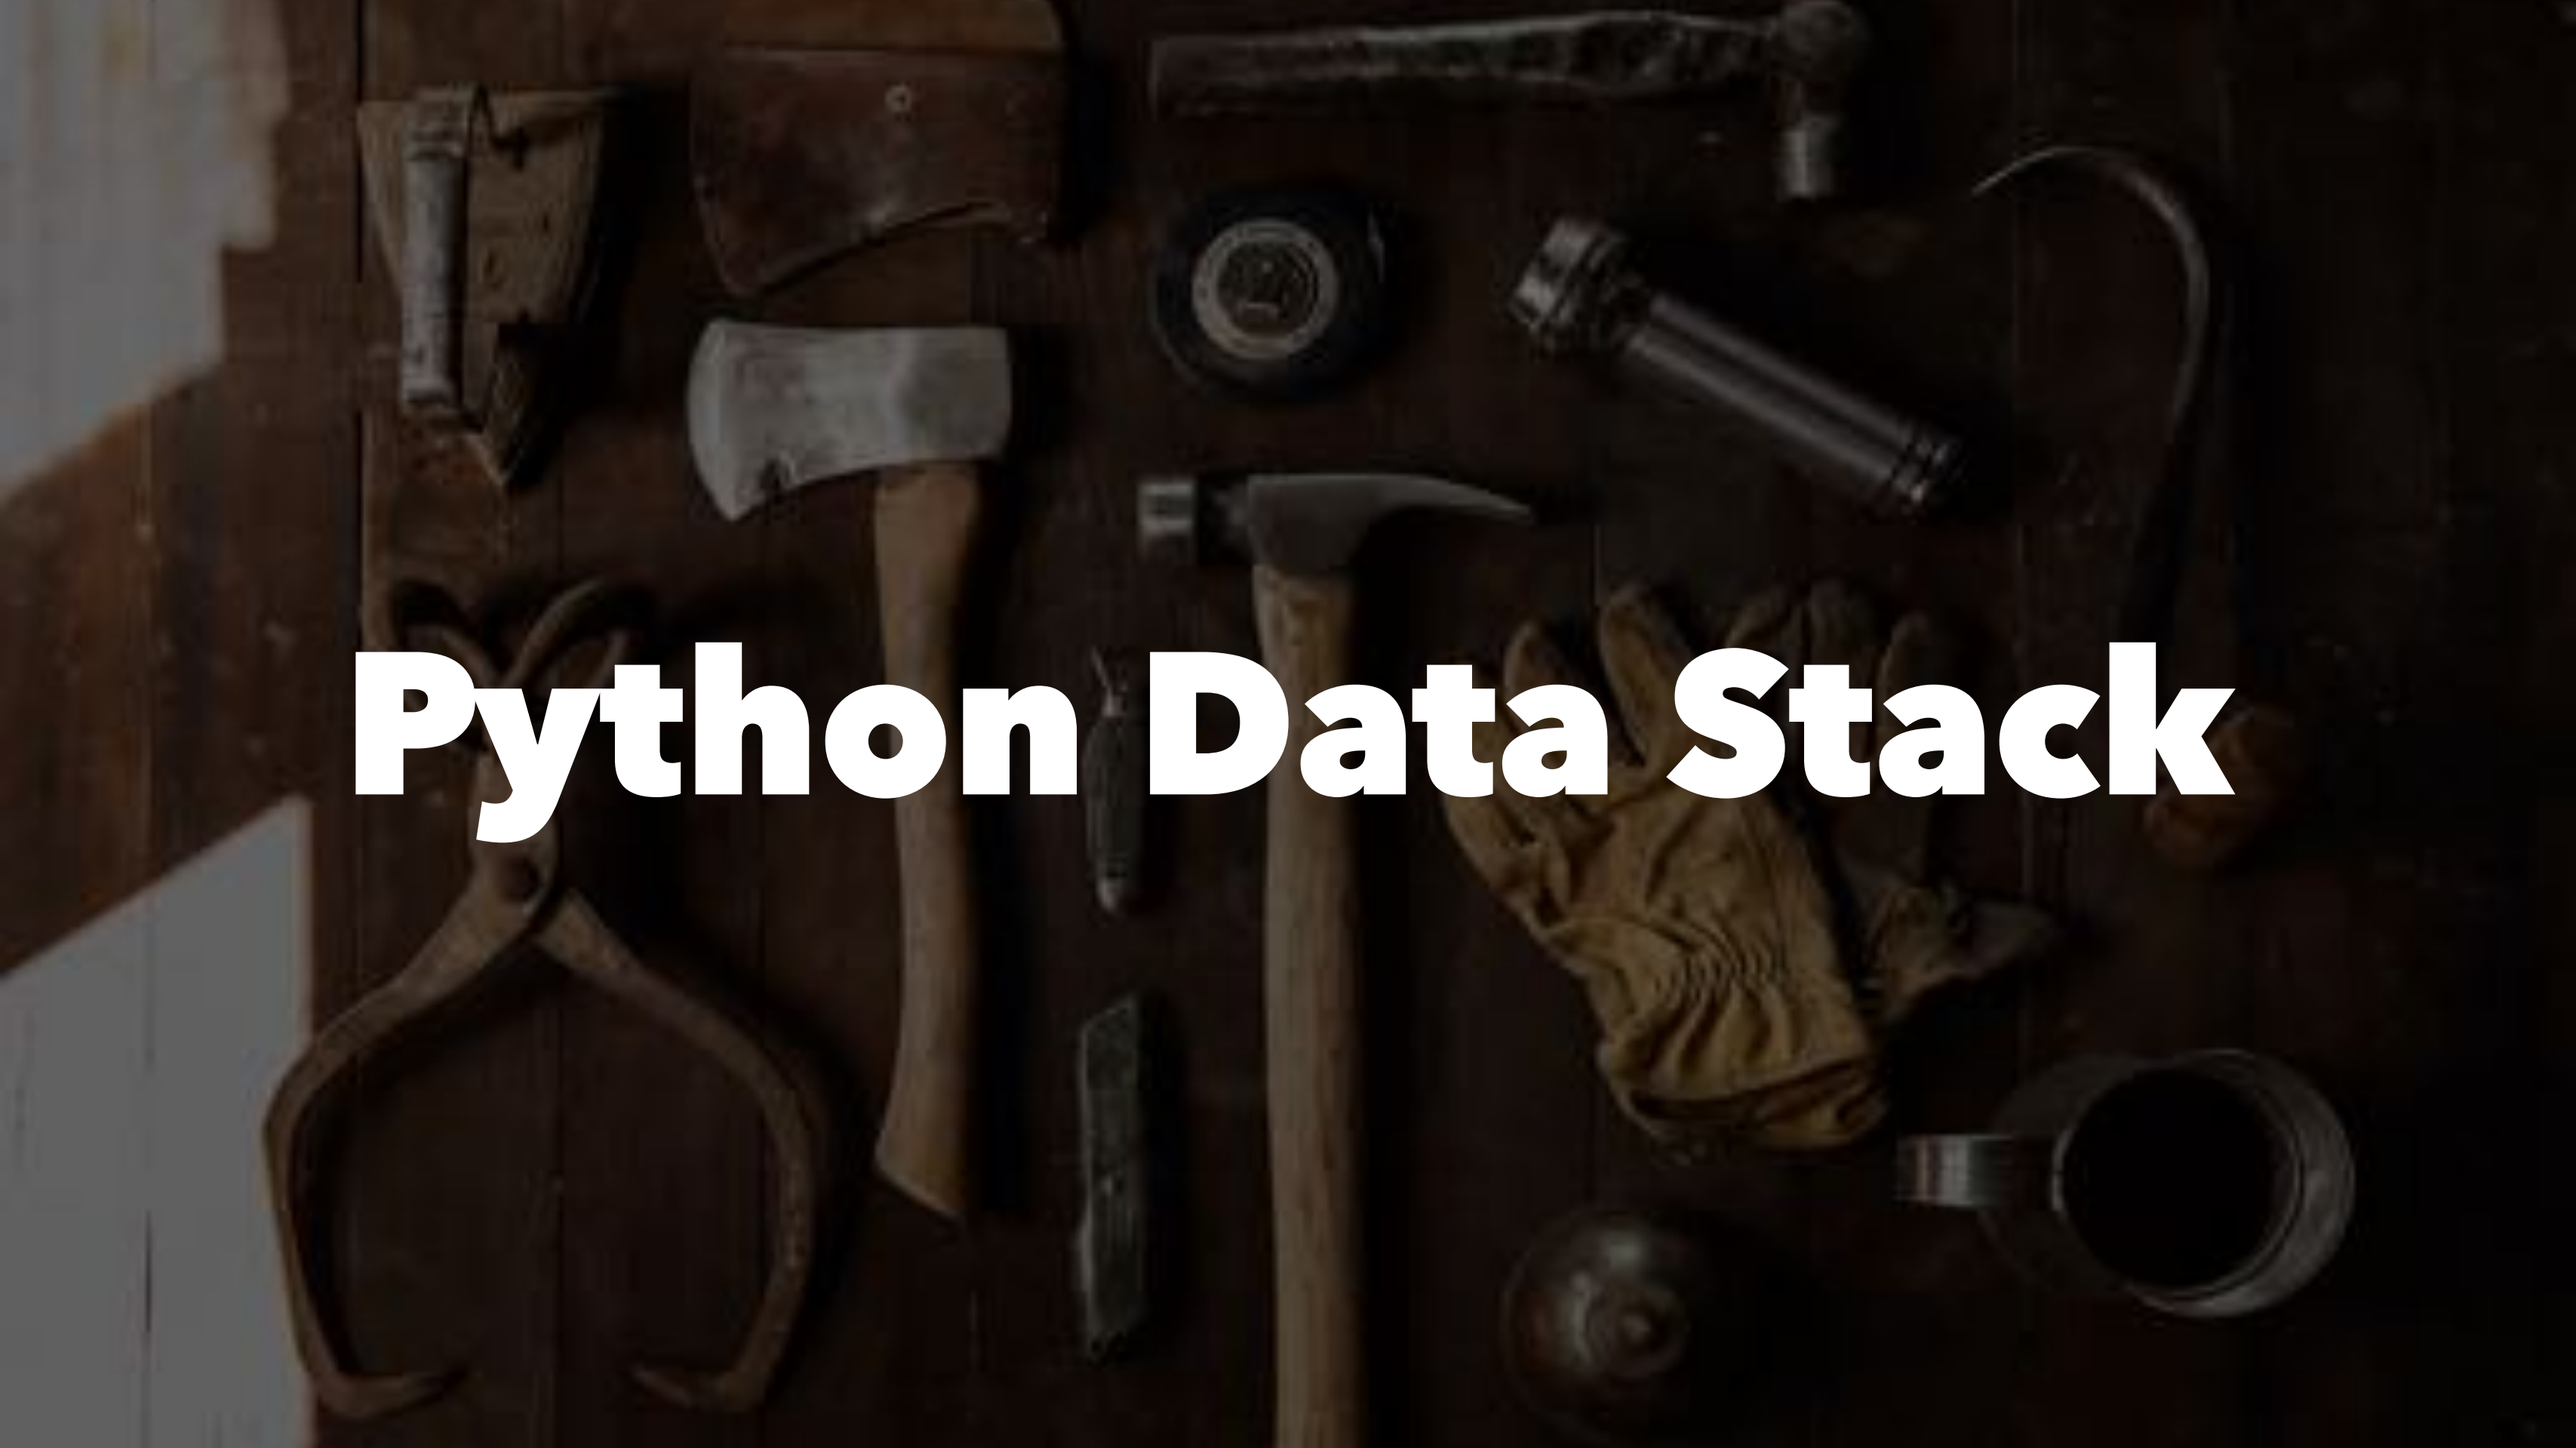
Python Data Stack


Case Studies

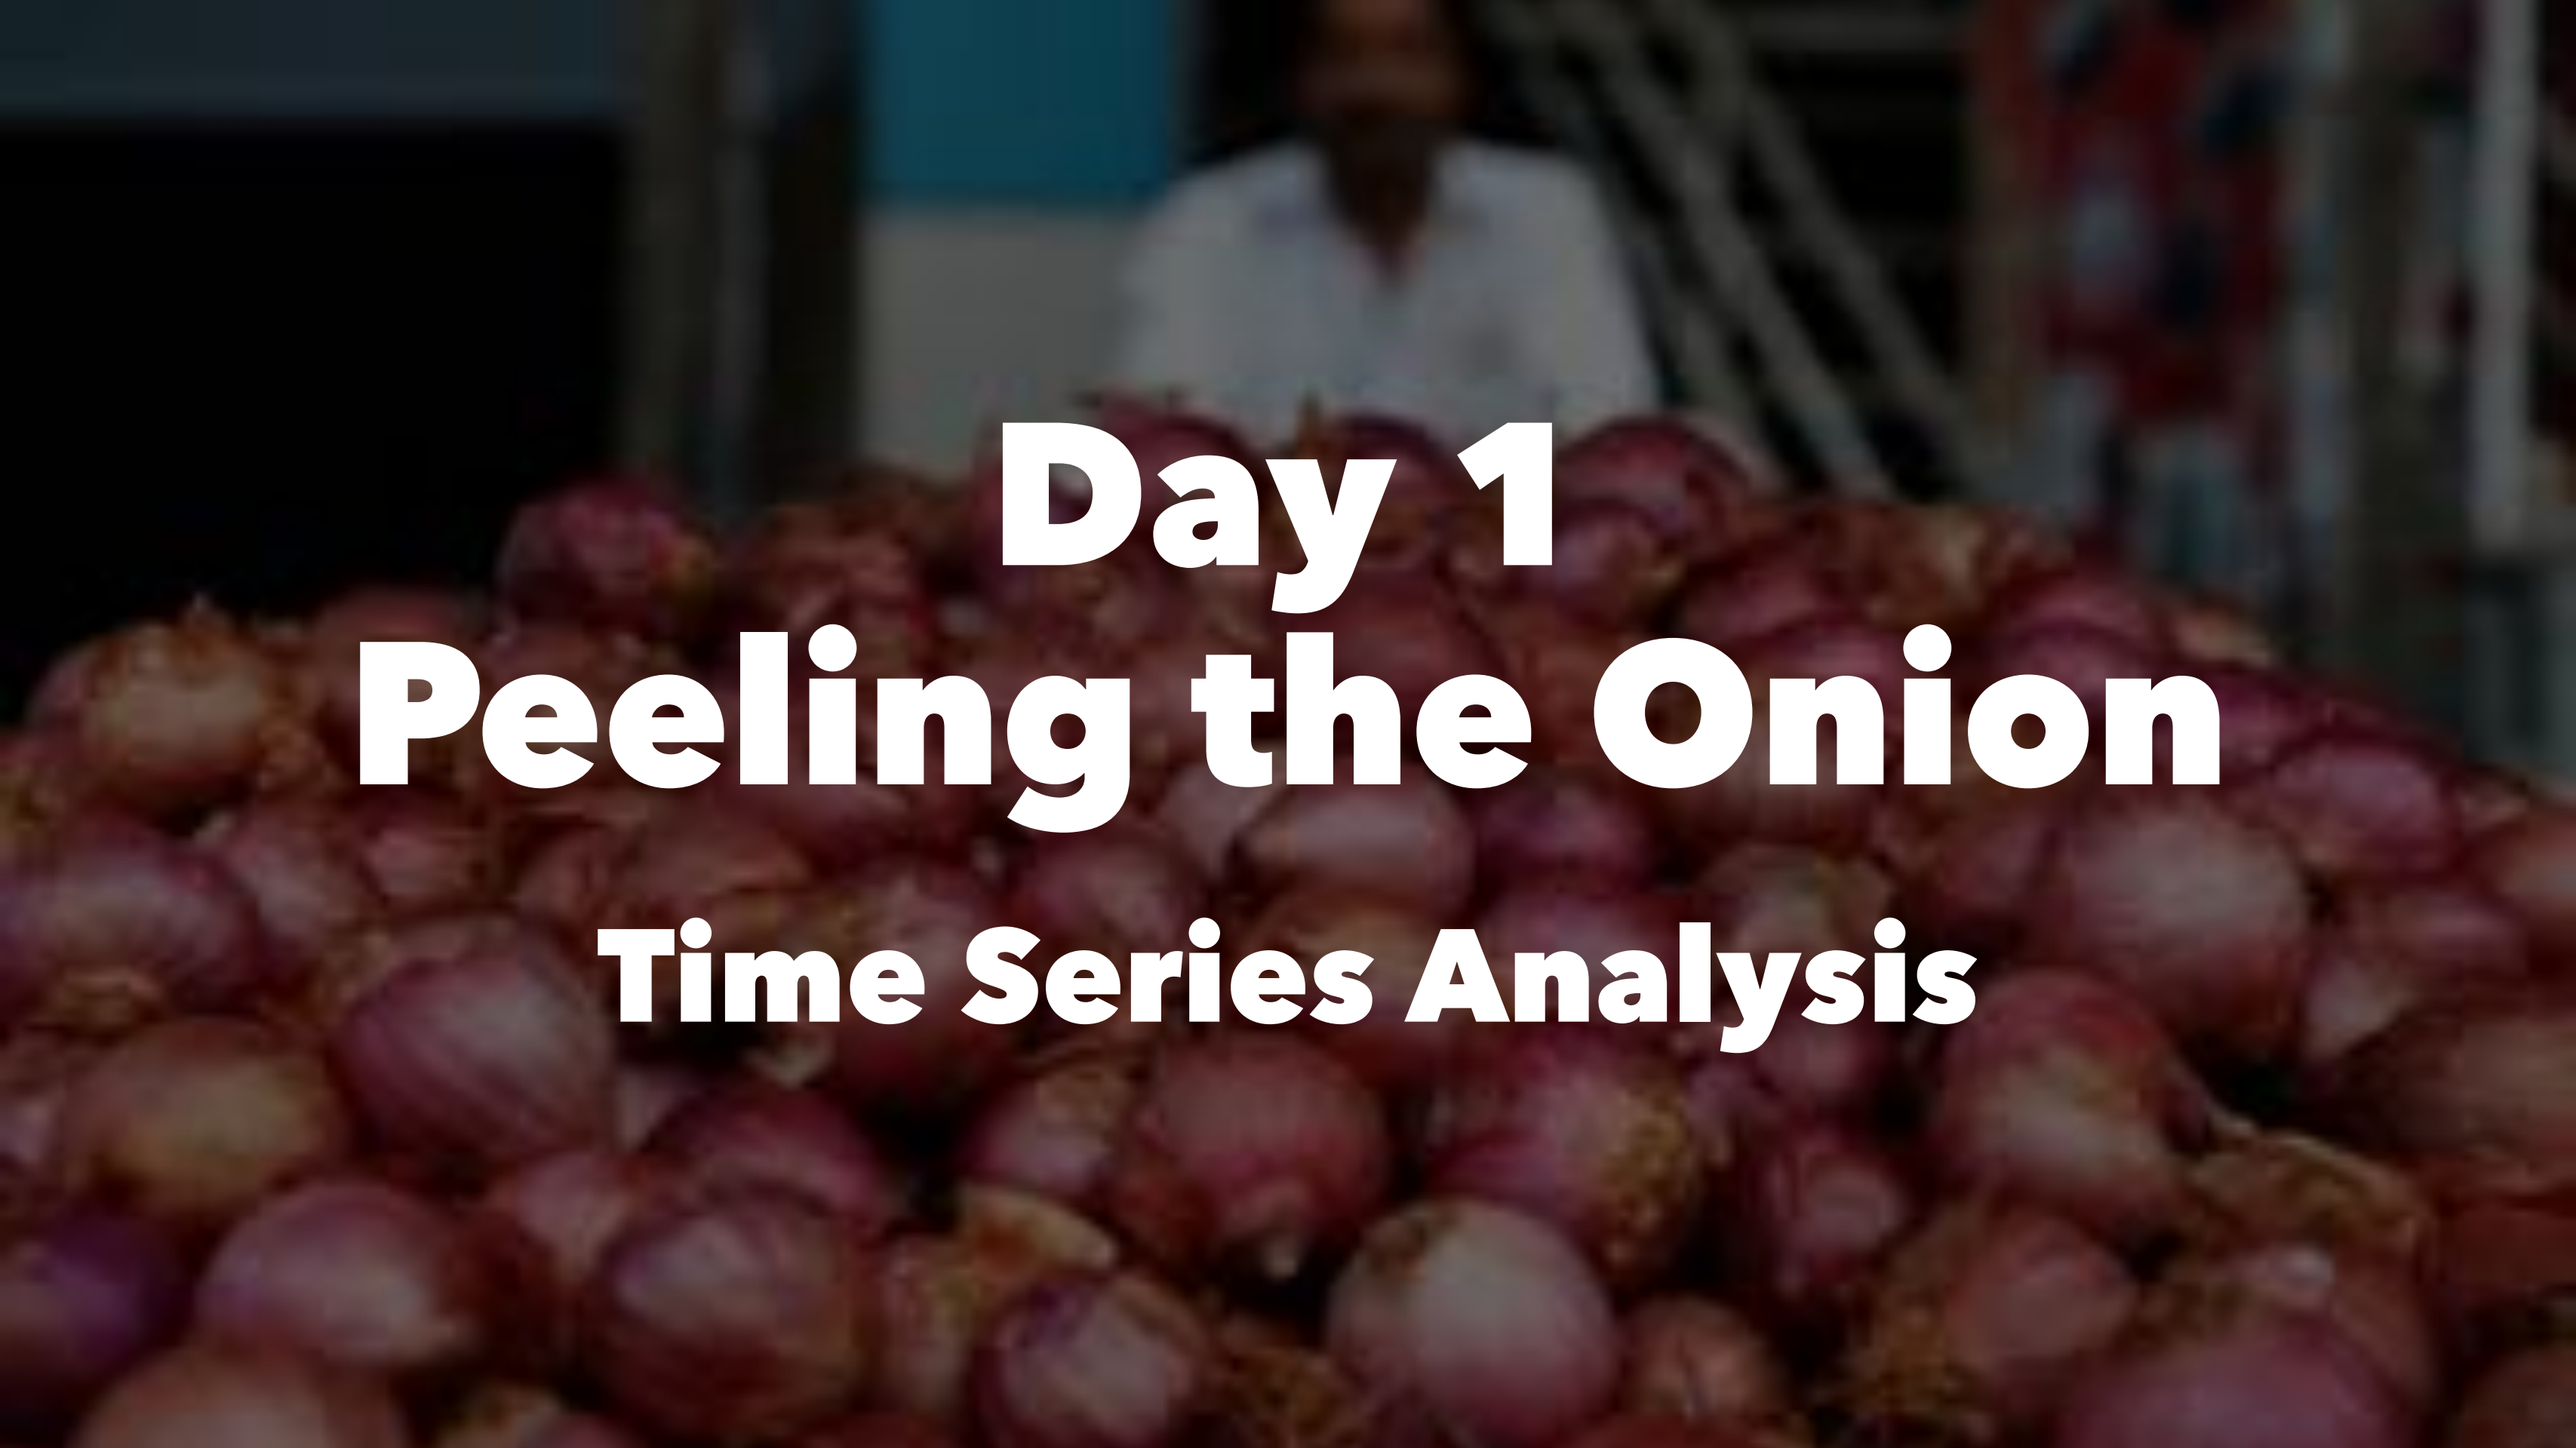

Day 1
Peeling the Onion
Time Series Analysis

Day 2
Grocery

Market Basket Analysis / Collaborative Filter

Day 2
BanK Marketing

Random Forest and Gradient Boosting

Day 3
DataTau

Text Analytics

Learning Approach


Do the Exercises


Pair up & Learn


Call for Help


Enjoy the workshop

Github Repo

https://github.com/amitkaps/machine-learning

Exercise

1. Time Series Exercise

"Predict the number of tickets that will be raised in the next week"

- **Frame**: What to forecast? At what horizon? At what level?
- **Acquire, Refine, Explore**: Do EDA to understand the trend and pattern within the data
- **Models**: Mean Model, Linear Trend, Random Walk, Simple Moving Average, Exp Smoothing, Decomposition, ARIMA
- **Insight**: Share the insight through a datavis of the models

2. Text Analytics Exercise

"Identify the entity, features and topics in the 'Comments Data' or "Twitter Data" within your text data"

- **Frame**: What are the comments you are trying to understand?
- **Acquire, Refine, Explore**: Do Wordcloud, Lemmatization, Part of Speech Analysis, and Entity Chunking
- **Models**: TF-IDF, Topic Modelling, Sentiment Analysis
- **Insight**: Share the insight through word cloud and topic

Feedback

https://amitkaps.typeform.com/to/i6wl2E

Recap

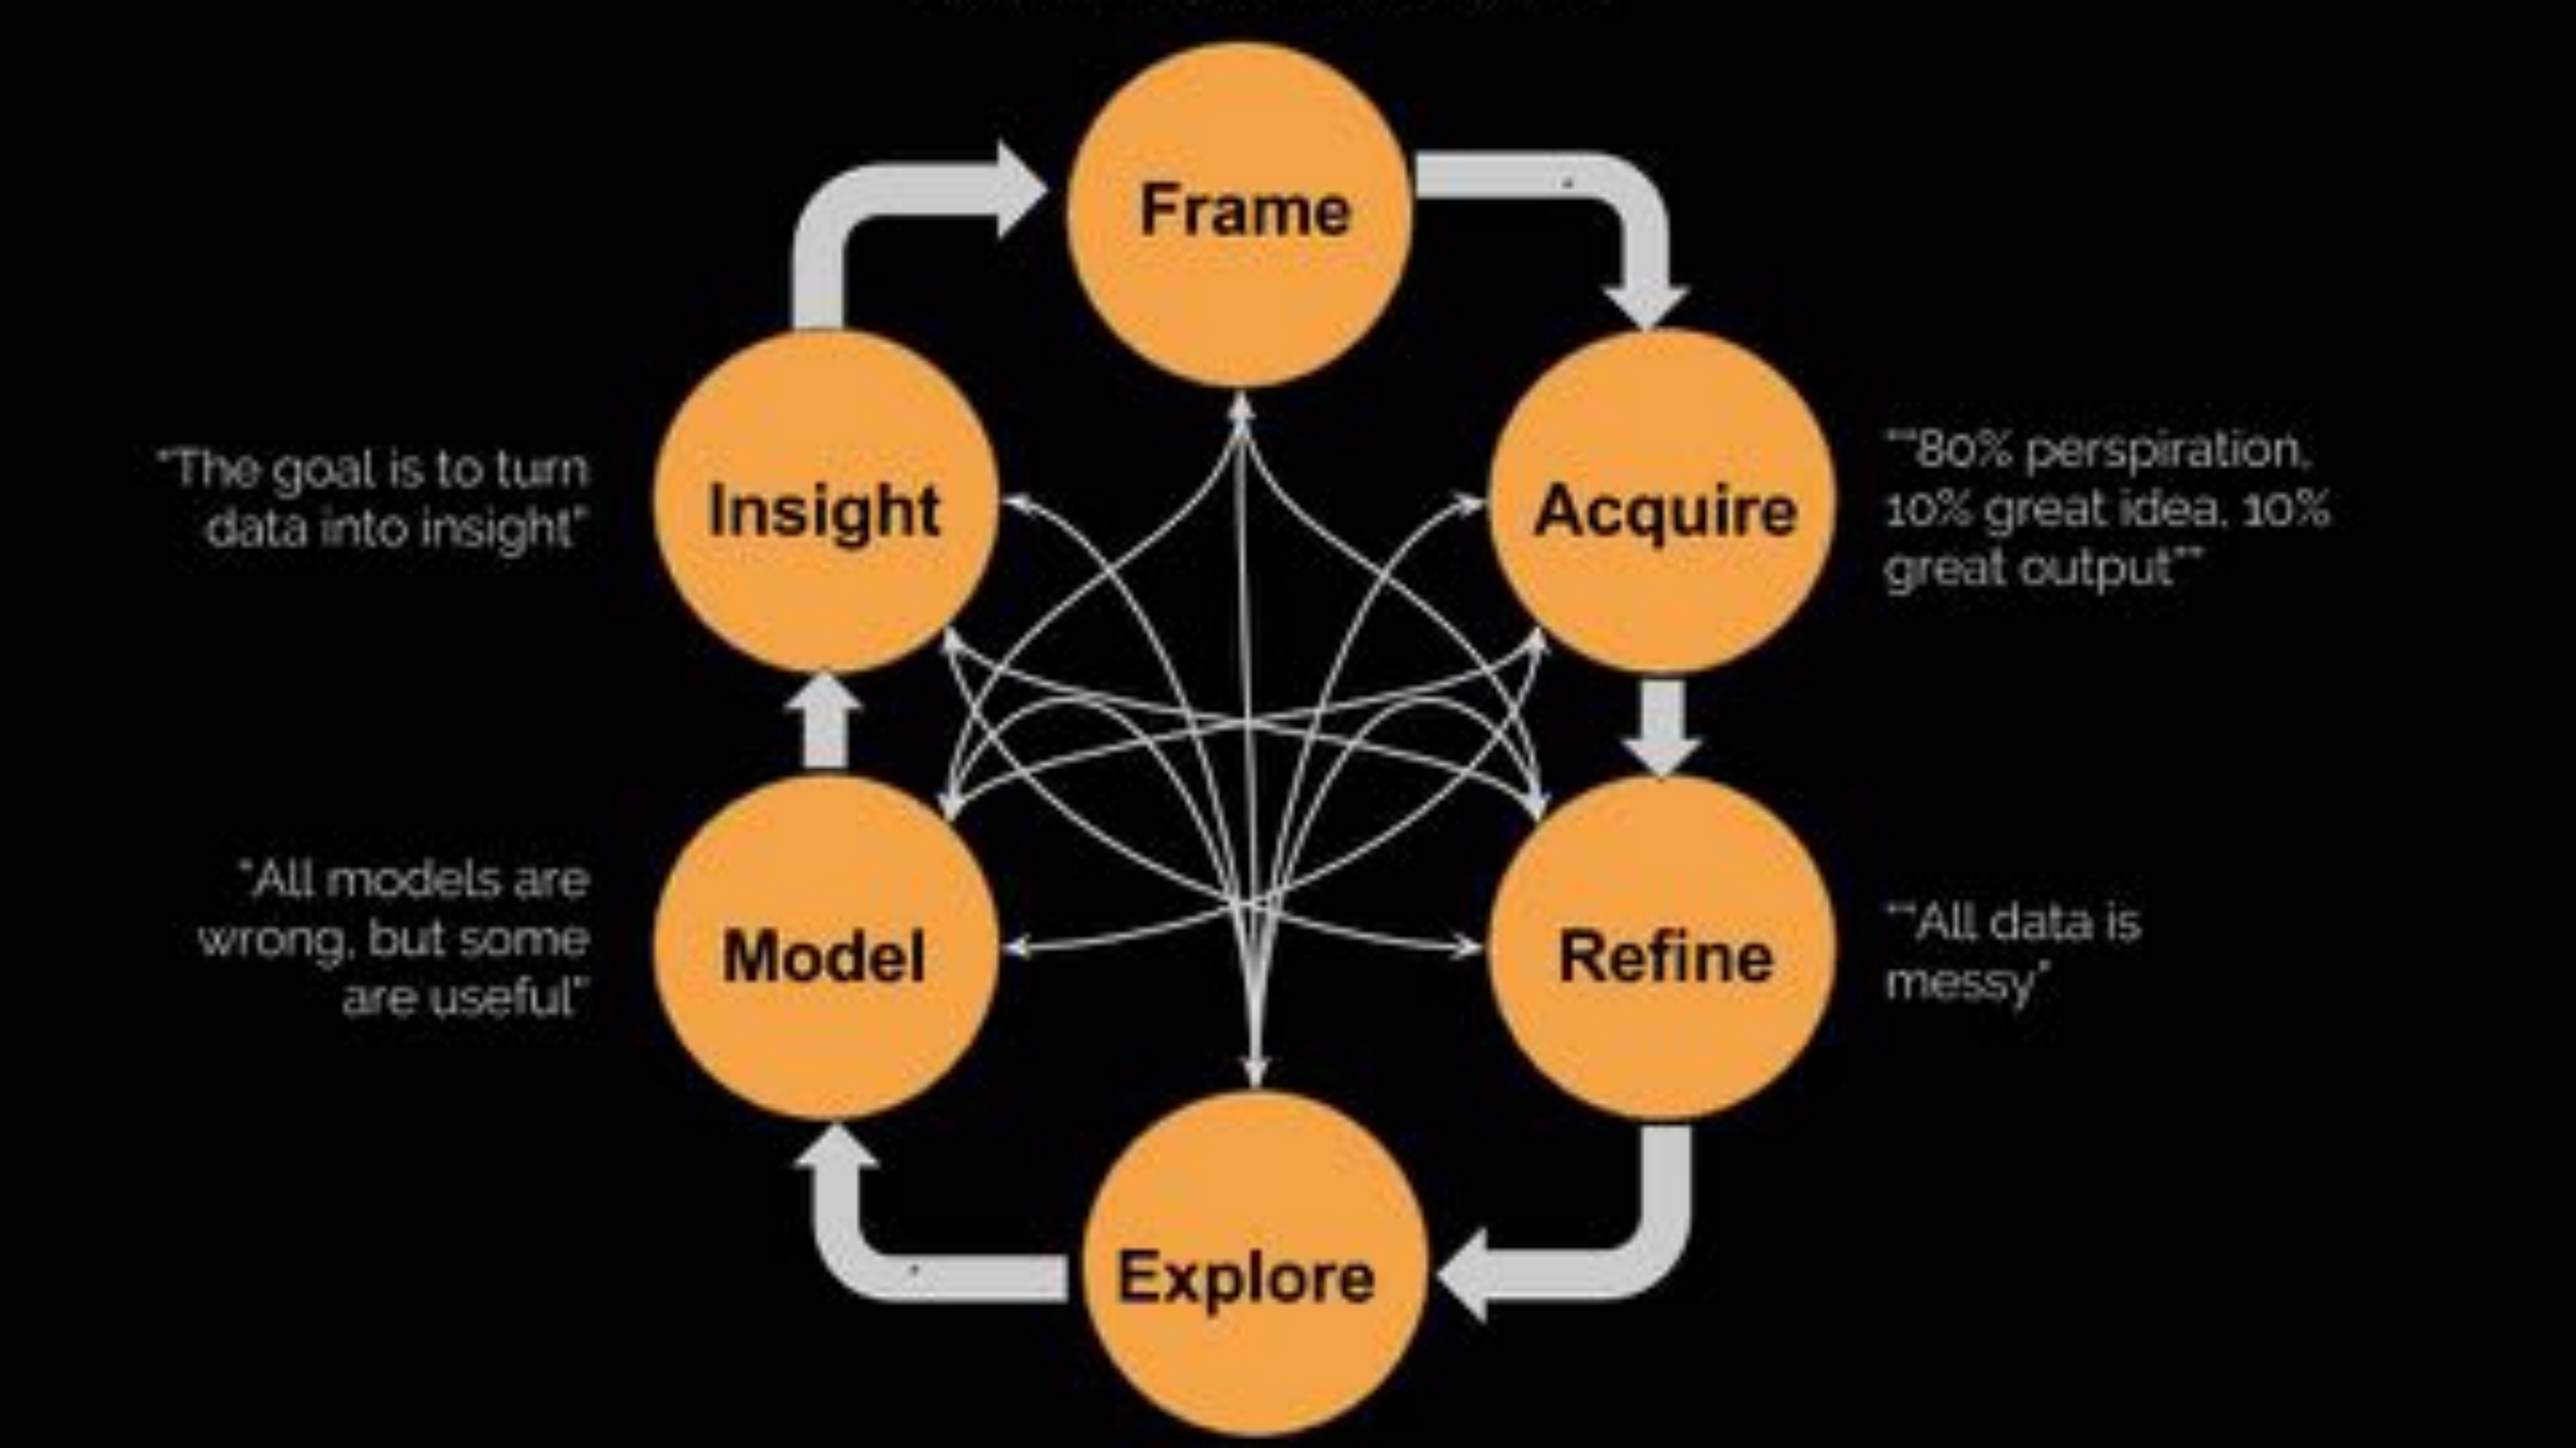
Frame
Acquire
Refine
Explore
Model
Insight
"The goal is to turn data into insight"
"80% perspiration, 10% great idea, 10% great output"
"All models are wrong, but some are useful"
"All data is messy"

Frame

- **Toy Problems**
- **Simple Problems**
- Complex Problems
- Business Problems
- Research Problems

Acquire

- **Scraping** (structured, unstructured)
- **Files** (csv, xls, json, xml, pdf, ...)
- Database (sqlite, ...)
- APIs
- Streaming

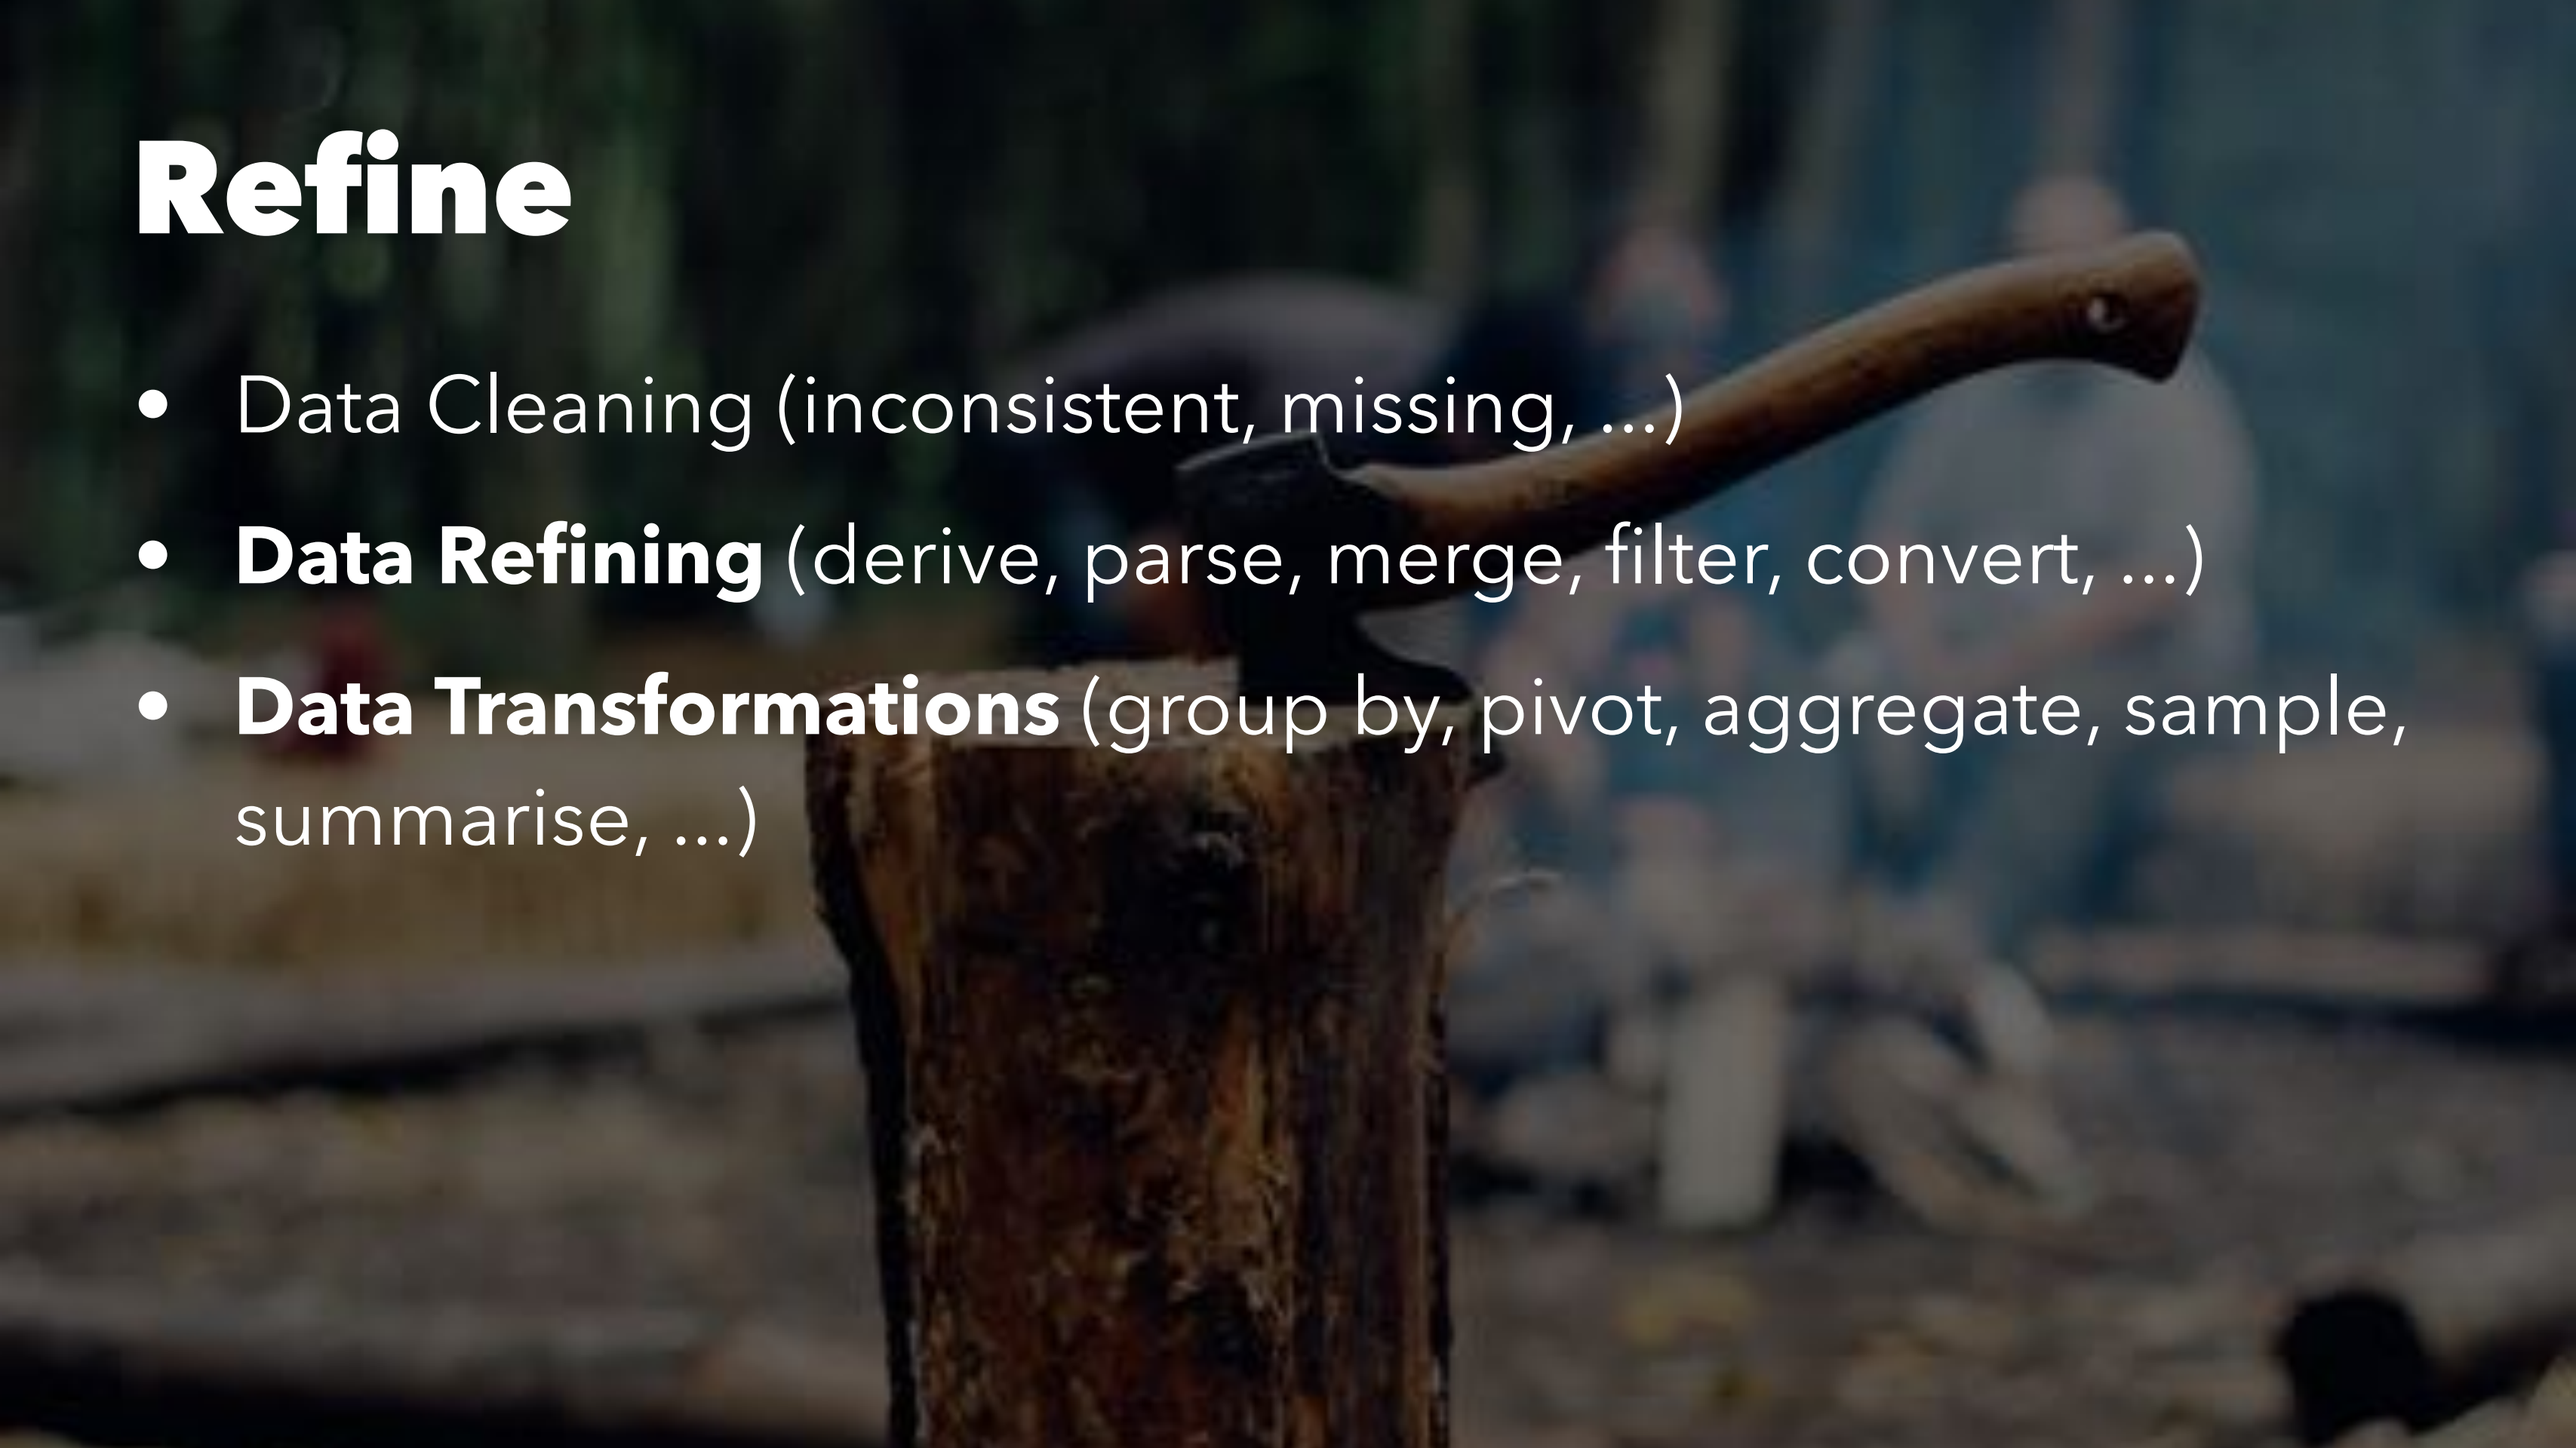

Refine

- Data Cleaning (inconsistent, missing, ...)
- **Data Refining** (derive, parse, merge, filter, convert, ...)
- **Data Transformations** (group by, pivot, aggregate, sample, summarise, ...)

Explore

- **Simple Vis**
- Multi Dimensional Vis
- Geographic Vis
- Large Data Vis (Bin - Summarise - Smooth)
- Interactive Vis

Model - Supervised Learning

- *Continuous*: Regression - **Linear**, Polynomial, Tree Based Methods - CART, **Random Forest**, Gradient Boosting Machines
- *Classification* - **Logistics Regression**, Tree, KNN, SVM, Naive-Bayes, Bayesian Network

Model - UnSupervised Learning

- *Continuous*: Clustering & Dimensionality Reduction like PCA, SVD, MDS, K-means
- *Categorical*: Association Analysis

Model - Advanced /

- **Time Series**
- **Text Analytics**
- Network / Graph Analytics
- Optimization

Model - Specialized

- Reinforcement Learning
- Online Learning
- Deep Learning
- Other Applications: Image, Speech

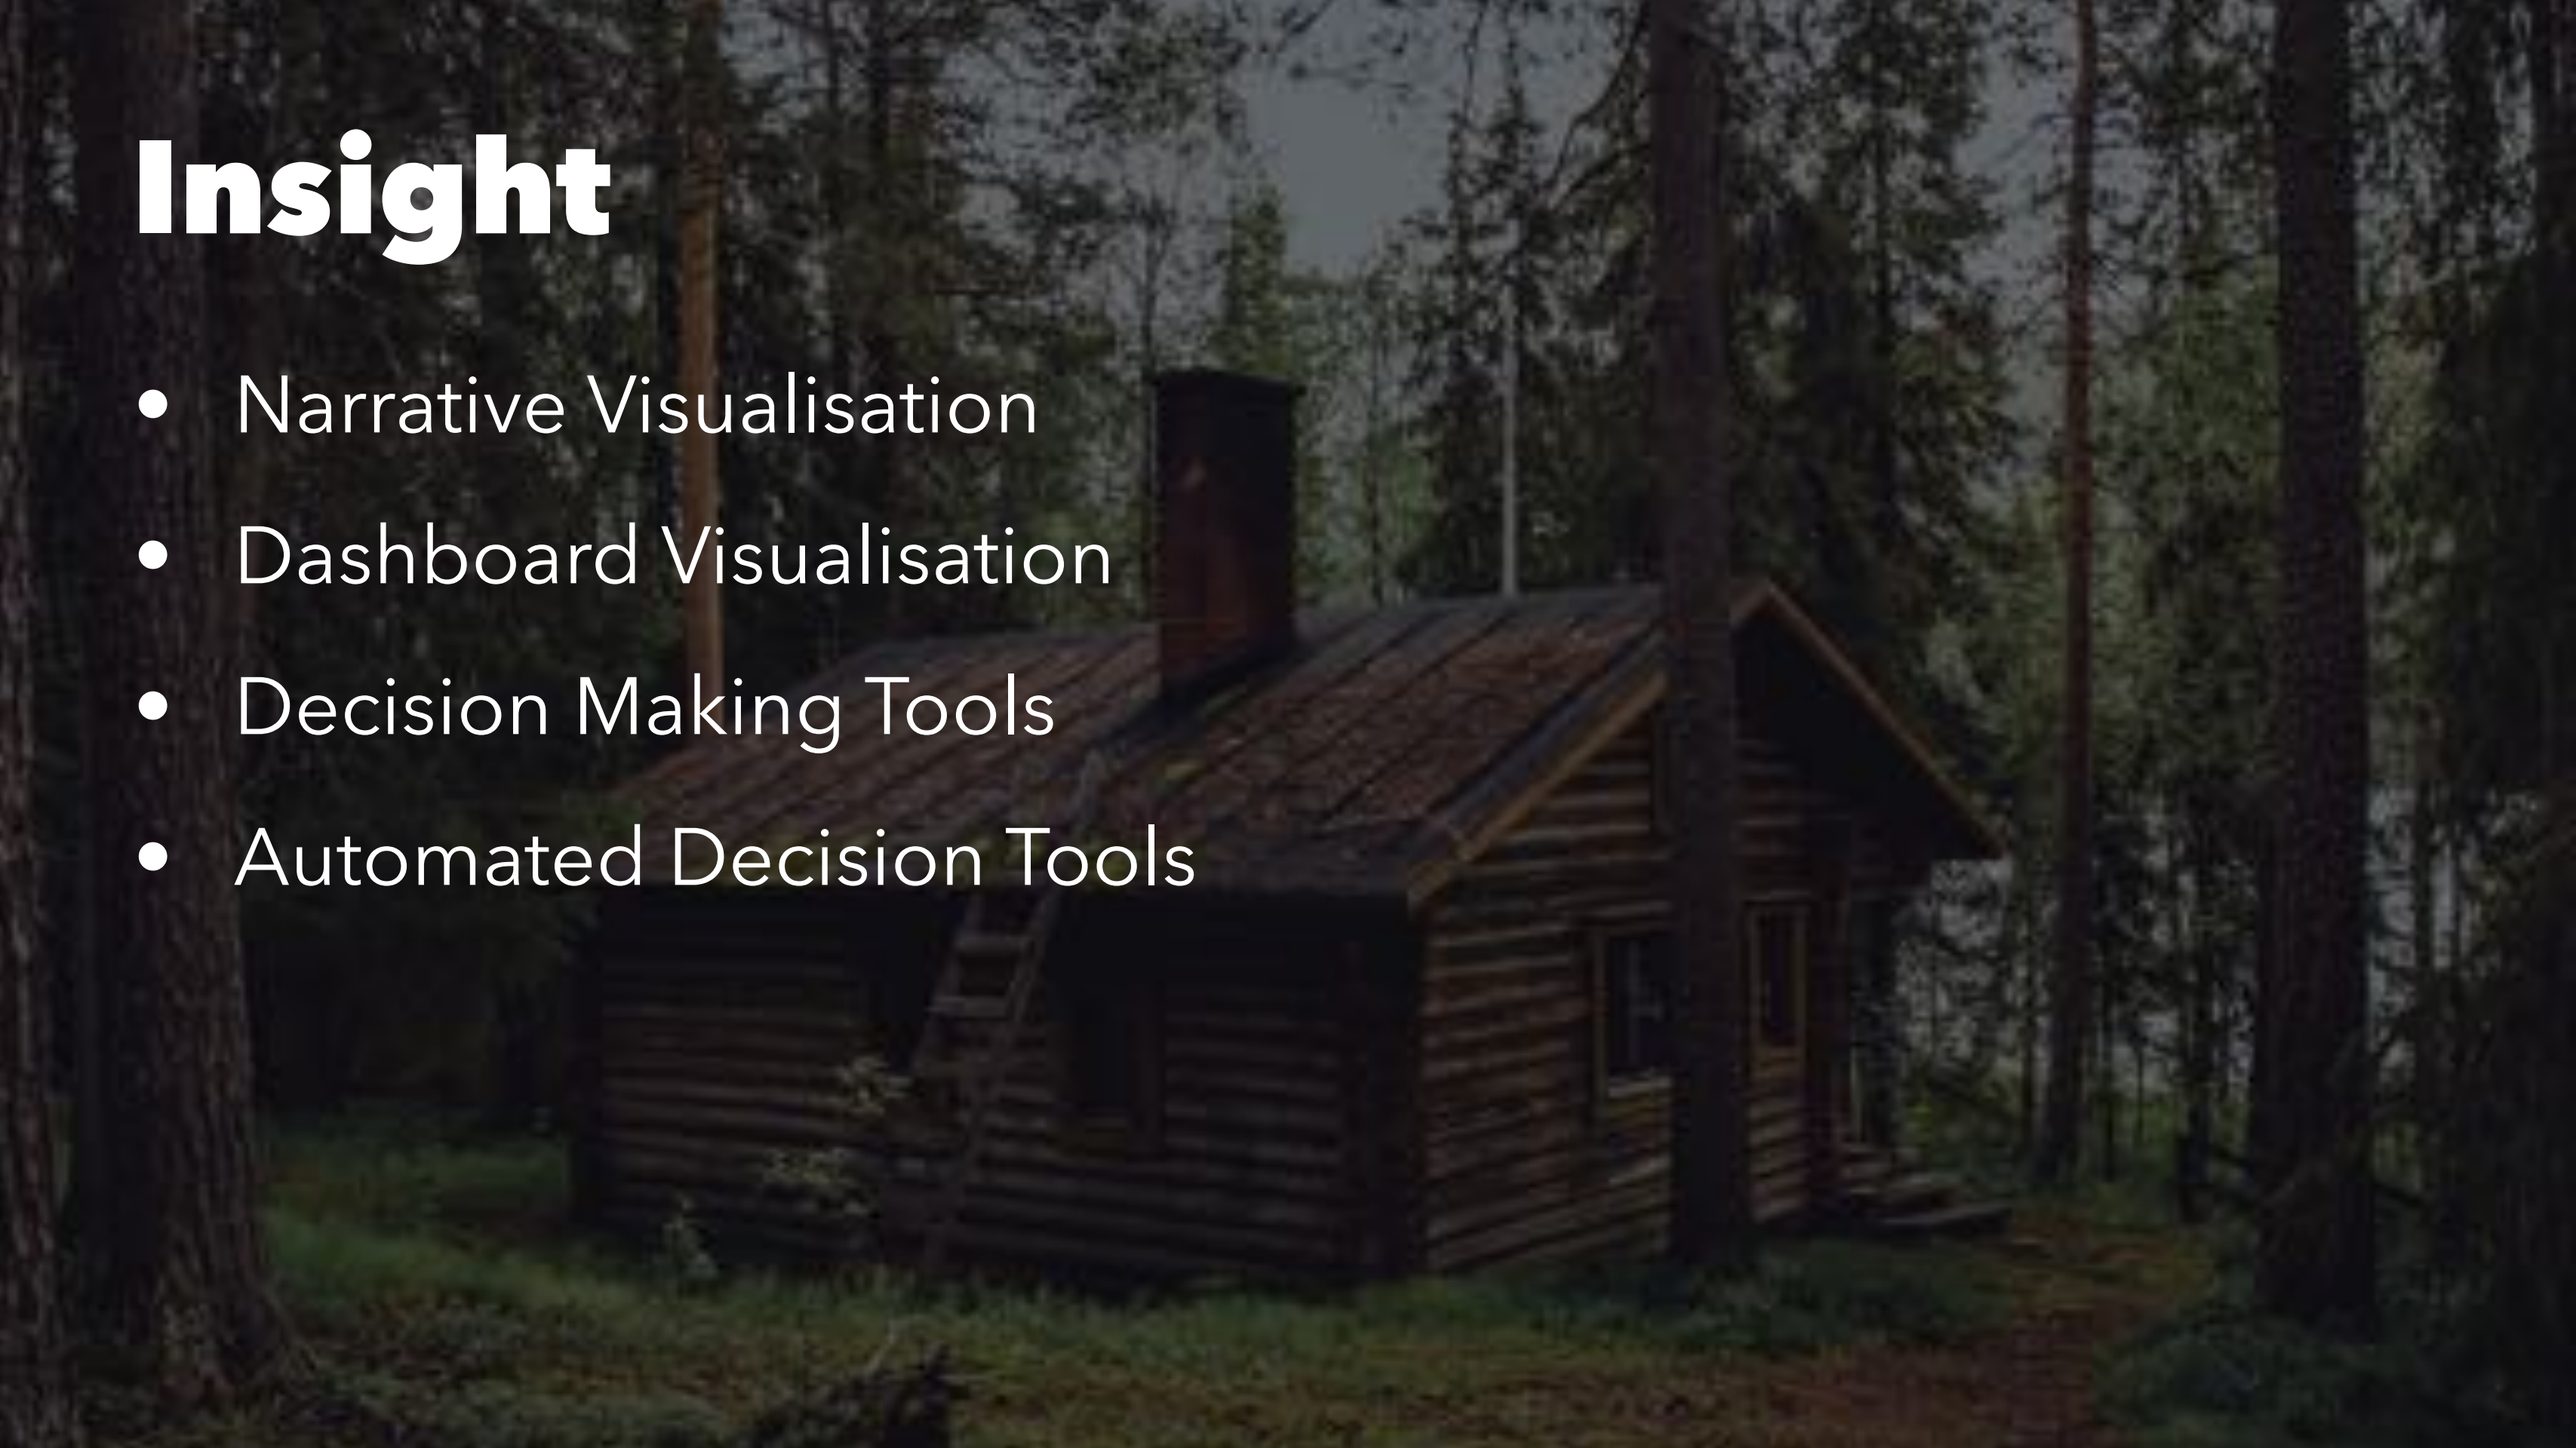

Insight

- Narrative Visualisation
- Dashboard Visualisation
- Decision Making Tools
- Automated Decision Tools

PyData Stack

- **Acquire / Refine**: Pandas, Beautiful Soup, Selenium, Requests, SQL Alchemy, Numpy, Blaze
- **Explore**: MatPlotLib, Seaborn, Bokeh, Plotly, Vega, Folium
- **Model**: Scikit-Learn, StatsModels, SciPy, Gensim, Keras, Tensor Flow, PySpark
- **Insight**: Django, Flask

Skills

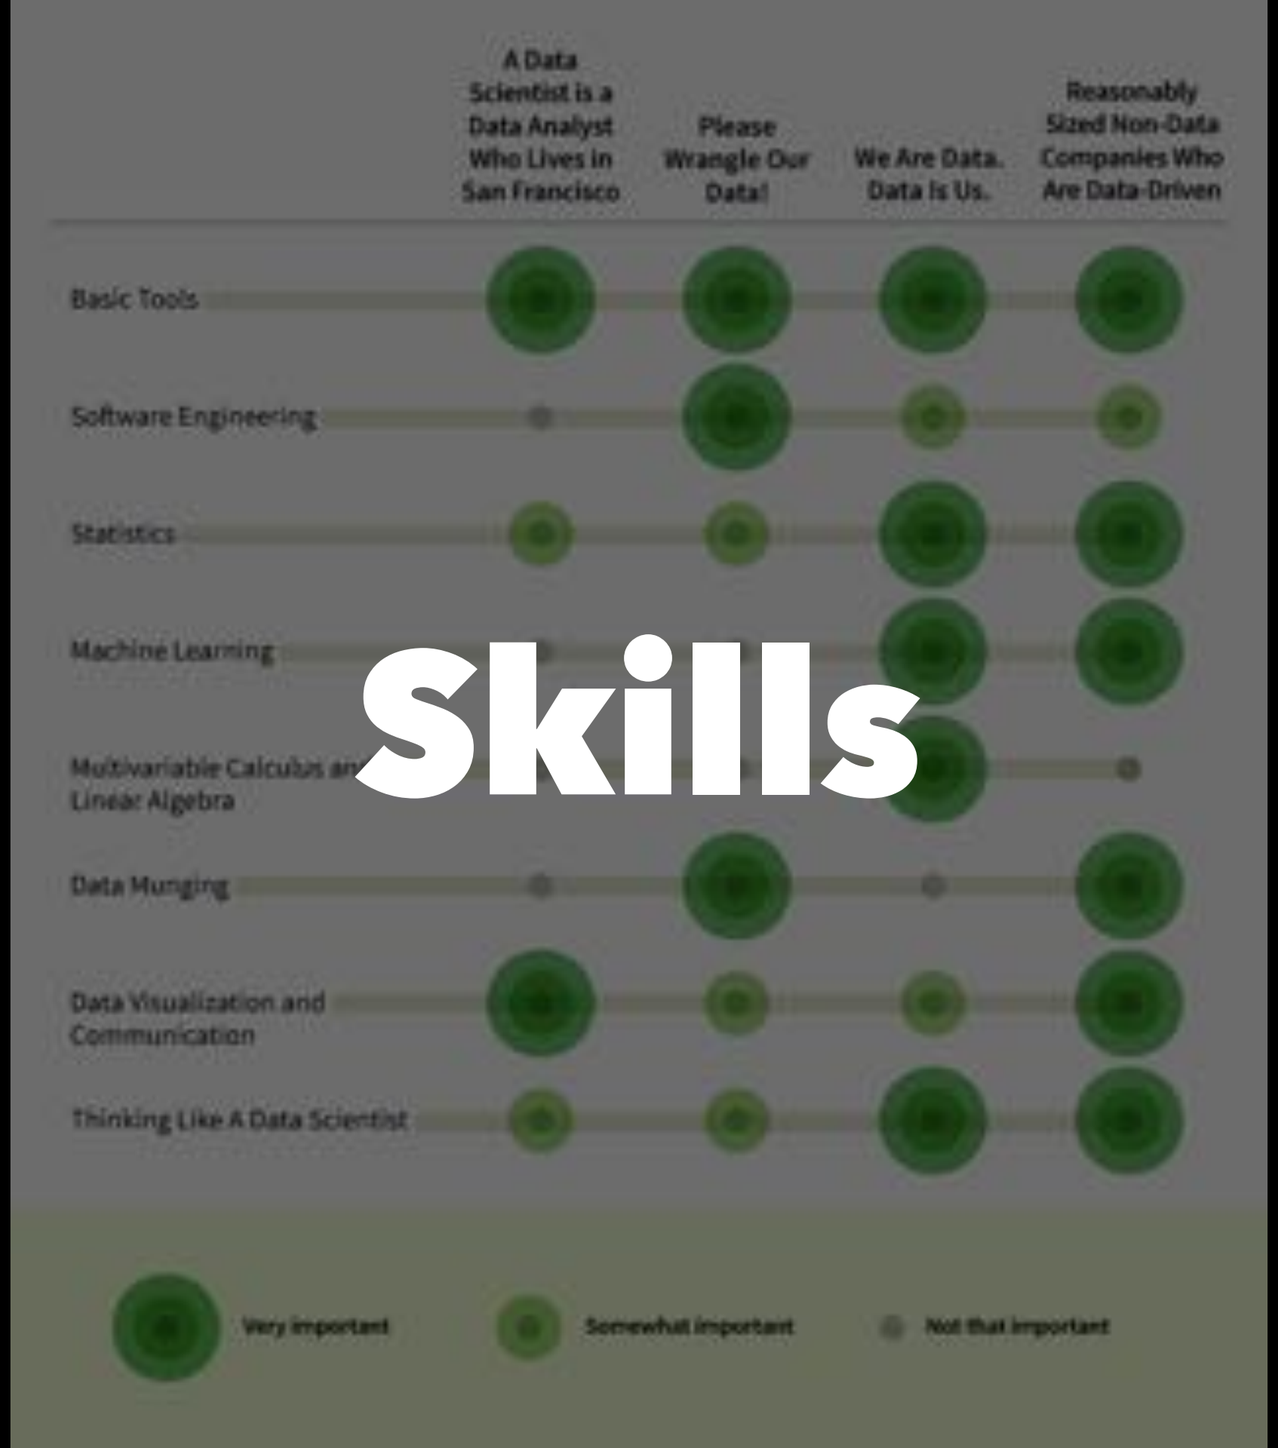
A Data Scientist is a Data Analyst Who Lives in San Francisco
Please Wrangle Our Data!
We Are Data. Data Is Us.
Reasonably Sized Non-Data Companies Who Are Data-Driven
Basic Tools
Software Engineering
Statistics
Machine Learning
Multivariable Calculus and Linear Algebra
Data Munging
Data Visualization and Communication
Thinking Like A Data Scientist
Very important
Somewhat important
Not that important

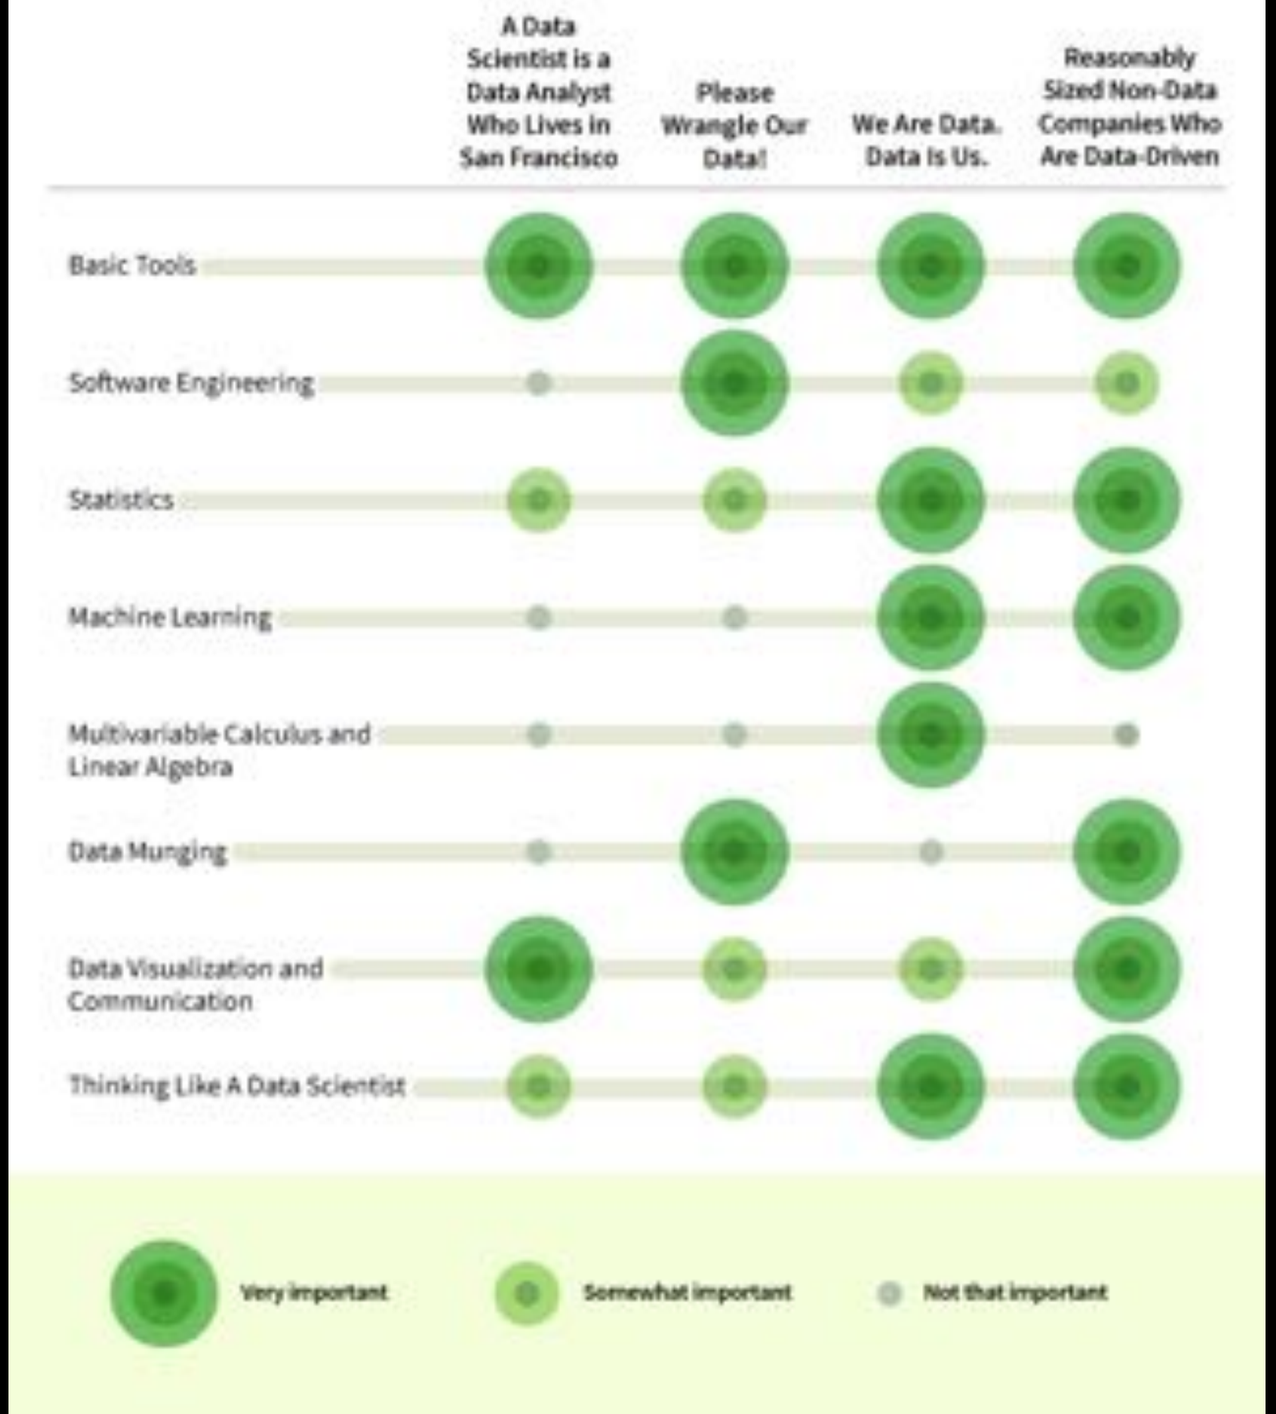
A Data Scientist is a Data Analyst Who Lives in San Francisco
Please Wrangle Our Data!
We Are Data. Data Is Us.
Reasonably Sized Non-Data Companies Who Are Data-Driven
Basic Tools
Software Engineering
Statistics
Machine Learning
Multivariable Calculus and Linear Algebra
Data Munging
Data Visualization and Communication
Thinking Like A Data Scientist
Very important
Somewhat important
Not that important

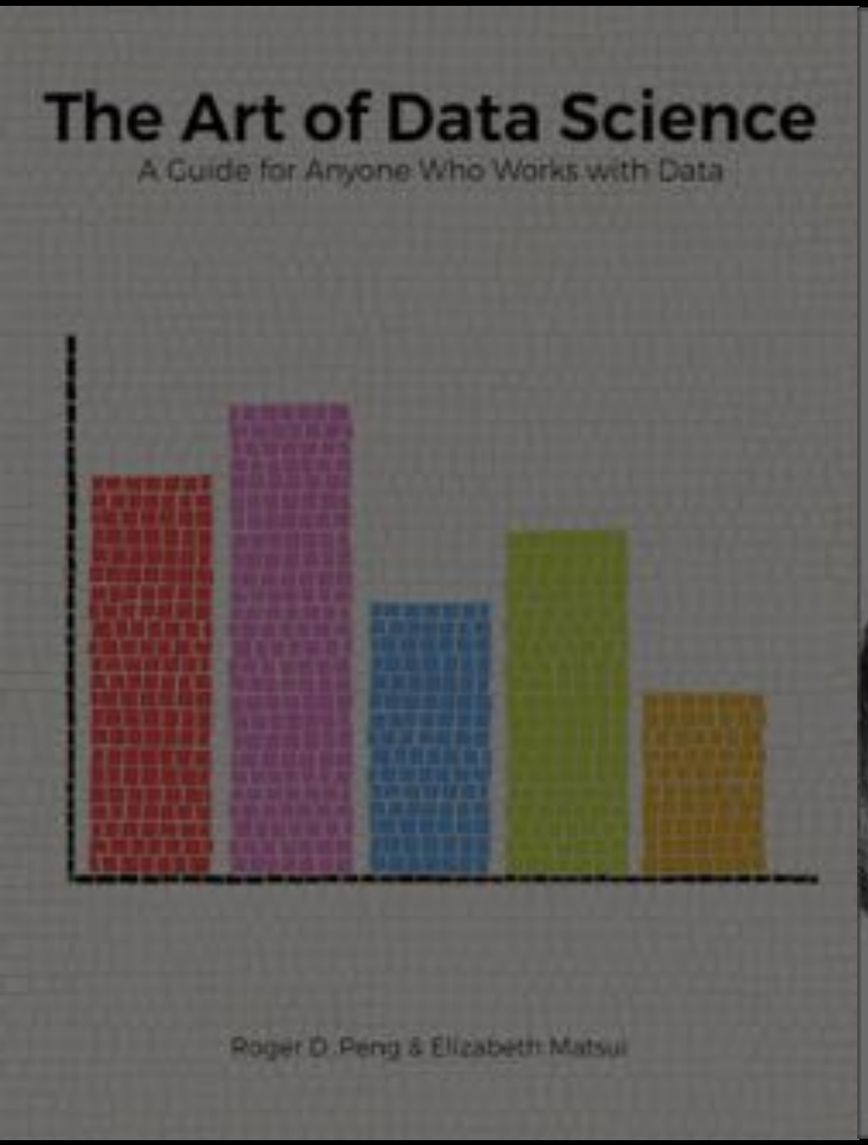
The Art of Data Science
A Guide for Anyone Who Works with Data
Roger D. Peng & Elizabeth Matsui

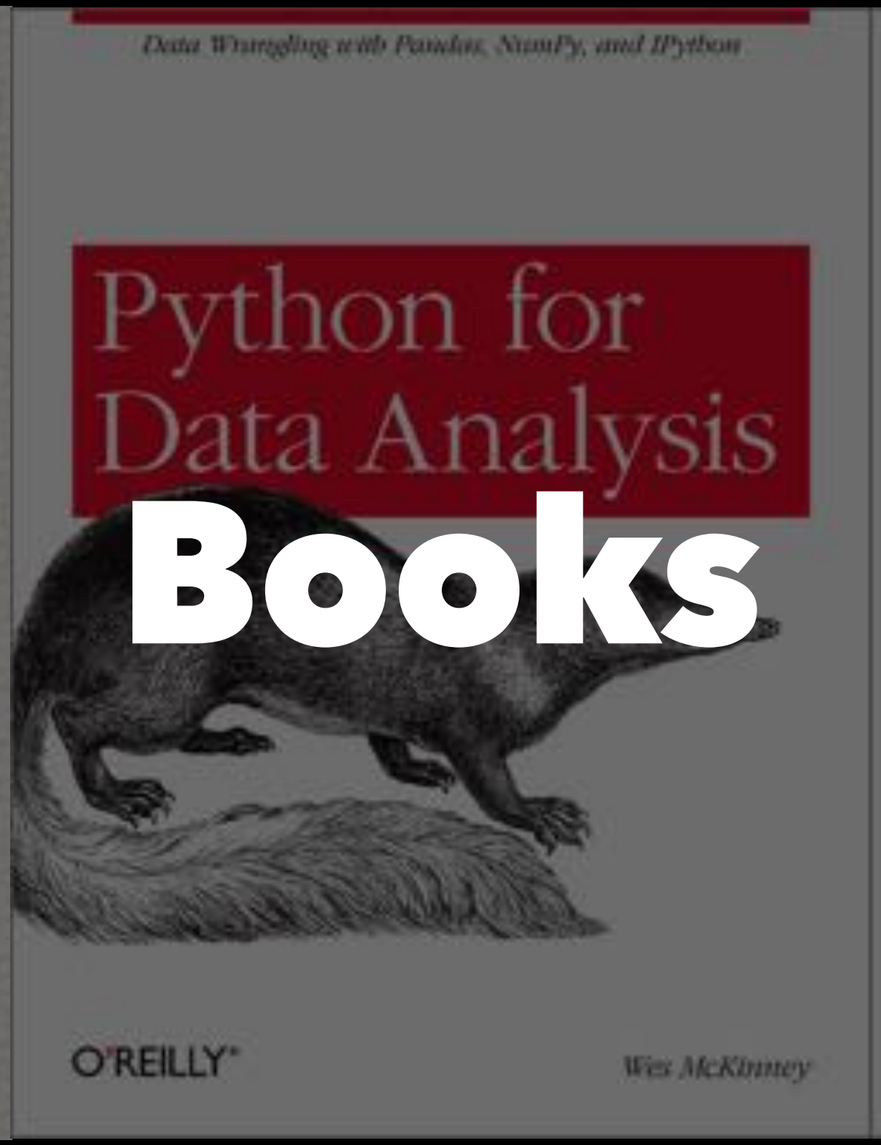
Data Wrangling with Pandas, NumPy, and IPython
Python for Data Analysis
O'REILLY
Wes McKinney

Books

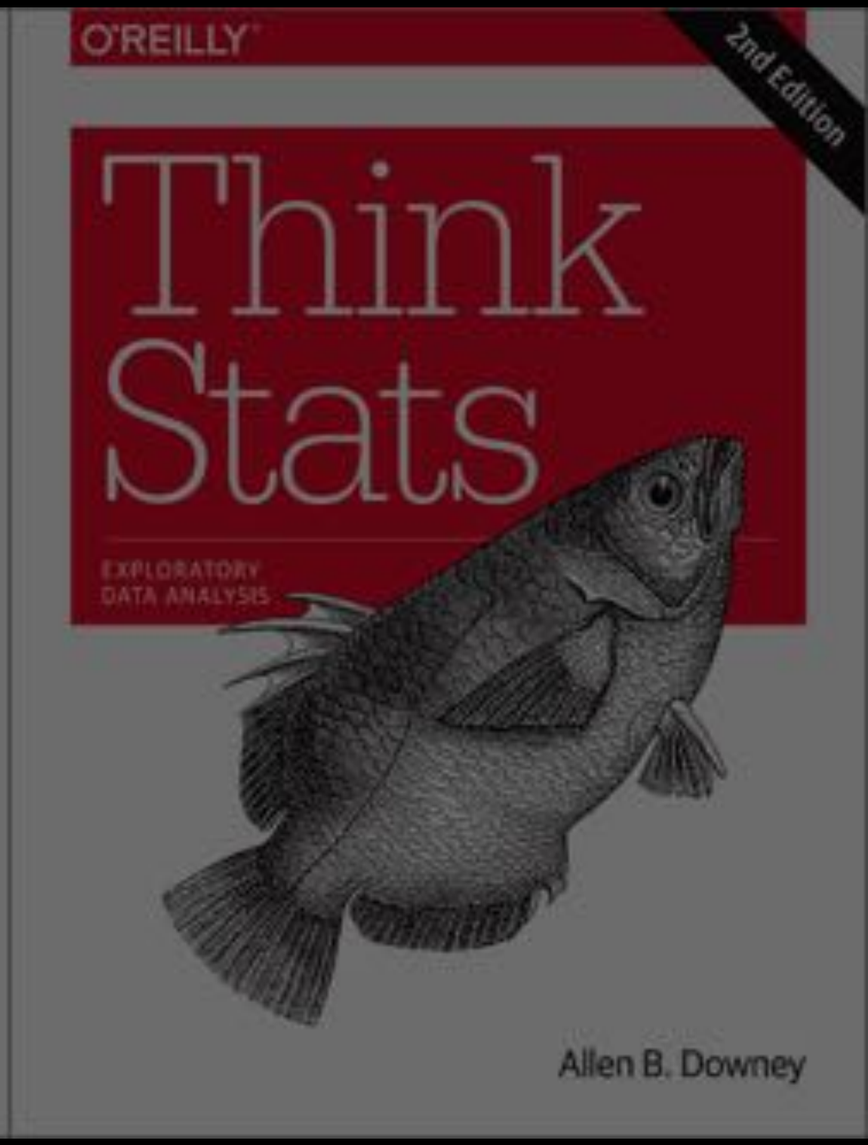
O'REILLY
2nd Edition
Think Stats
EXPLORATORY DATA ANALYSIS
Allen B. Downey

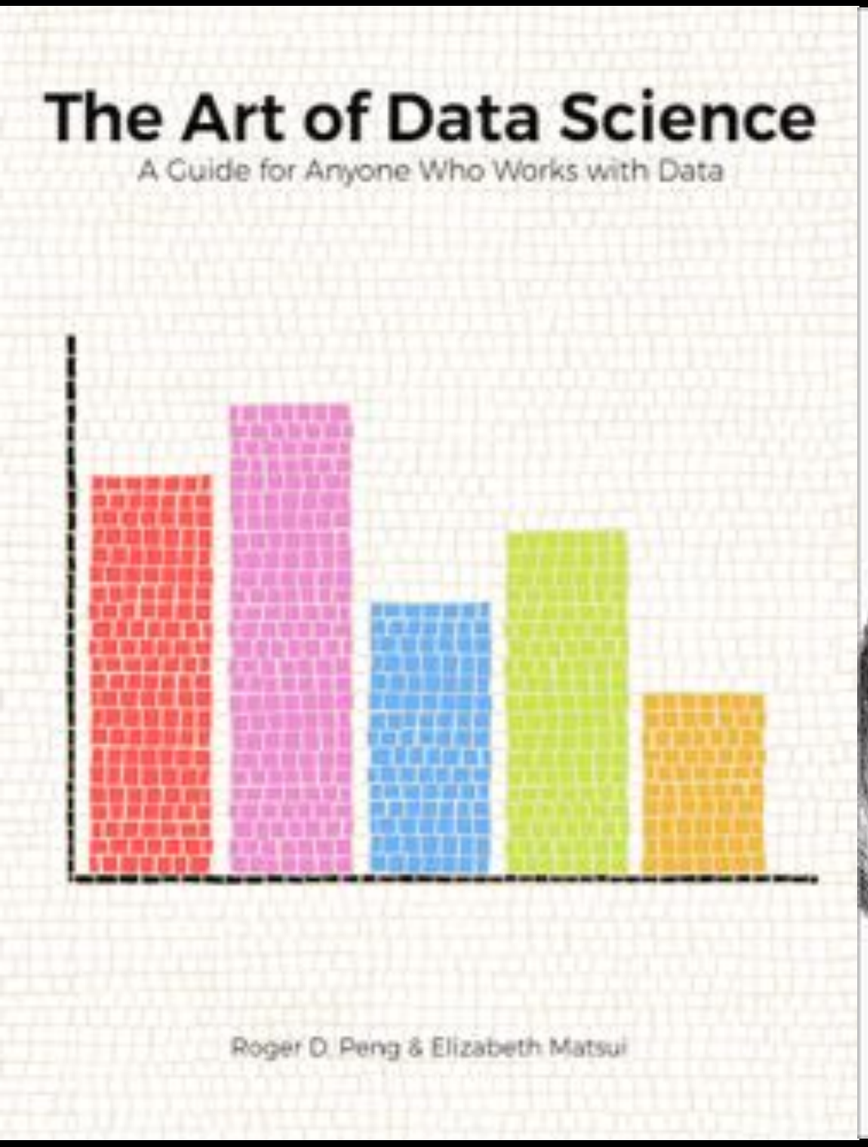
The Art of Data Science
A Guide for Anyone Who Works with Data
Roger D. Peng & Elizabeth Matsui

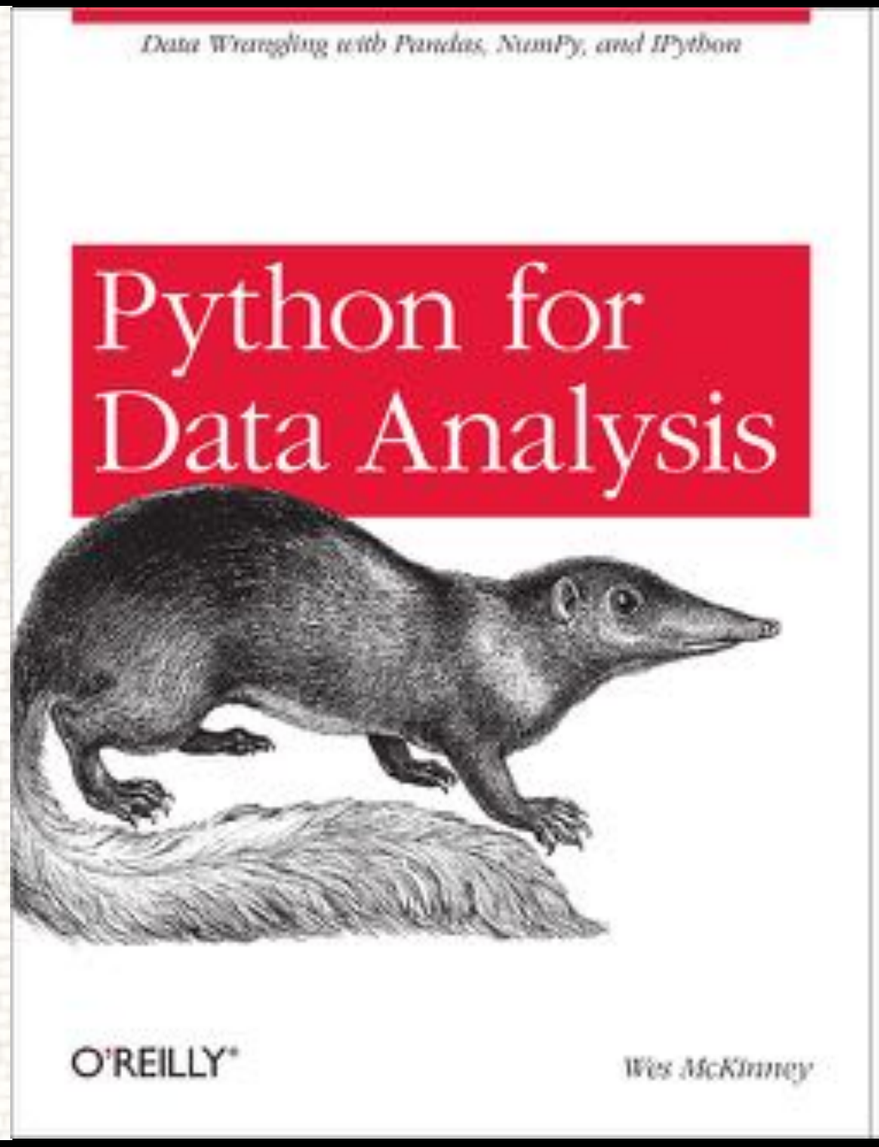
Data Wrangling with Pandas, NumPy, and IPython
Python for Data Analysis
O'REILLY
Wes McKinney

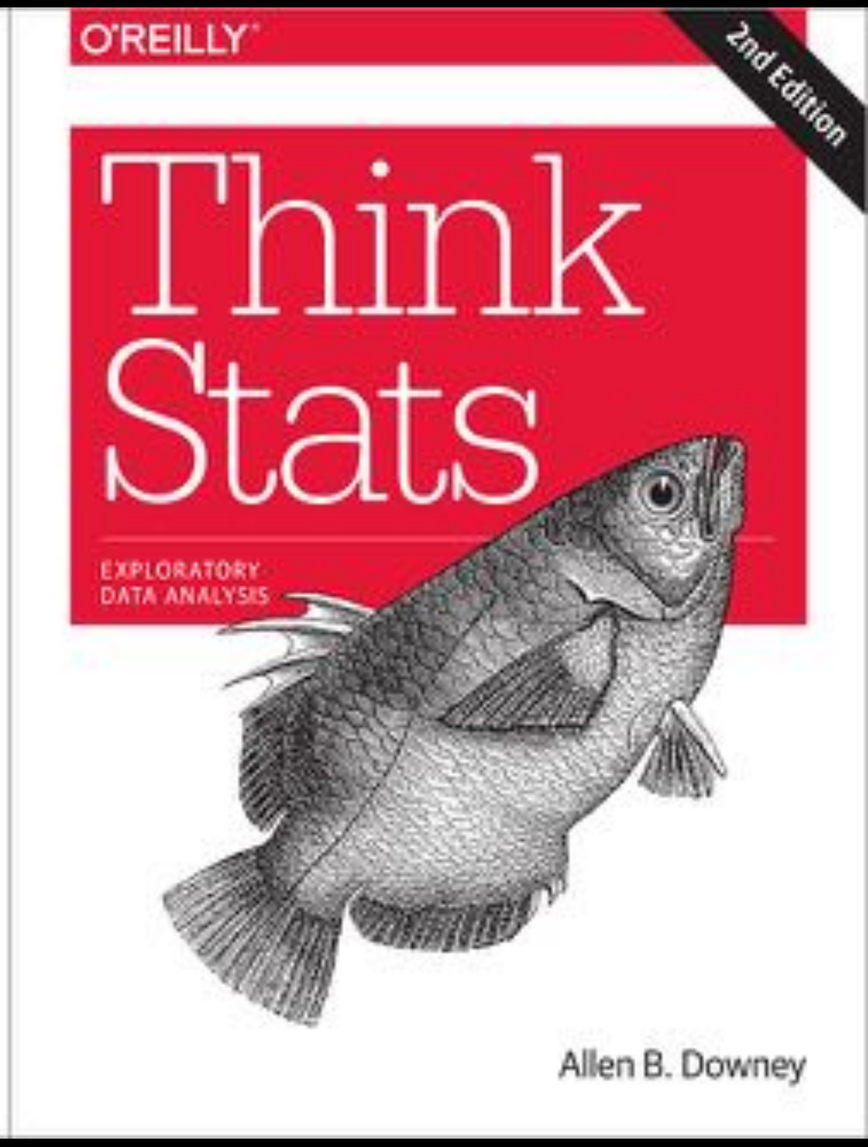
O'REILLY
2nd Edition
Think Stats
EXPLORATORY DATA ANALYSIS
Allen B. Downey

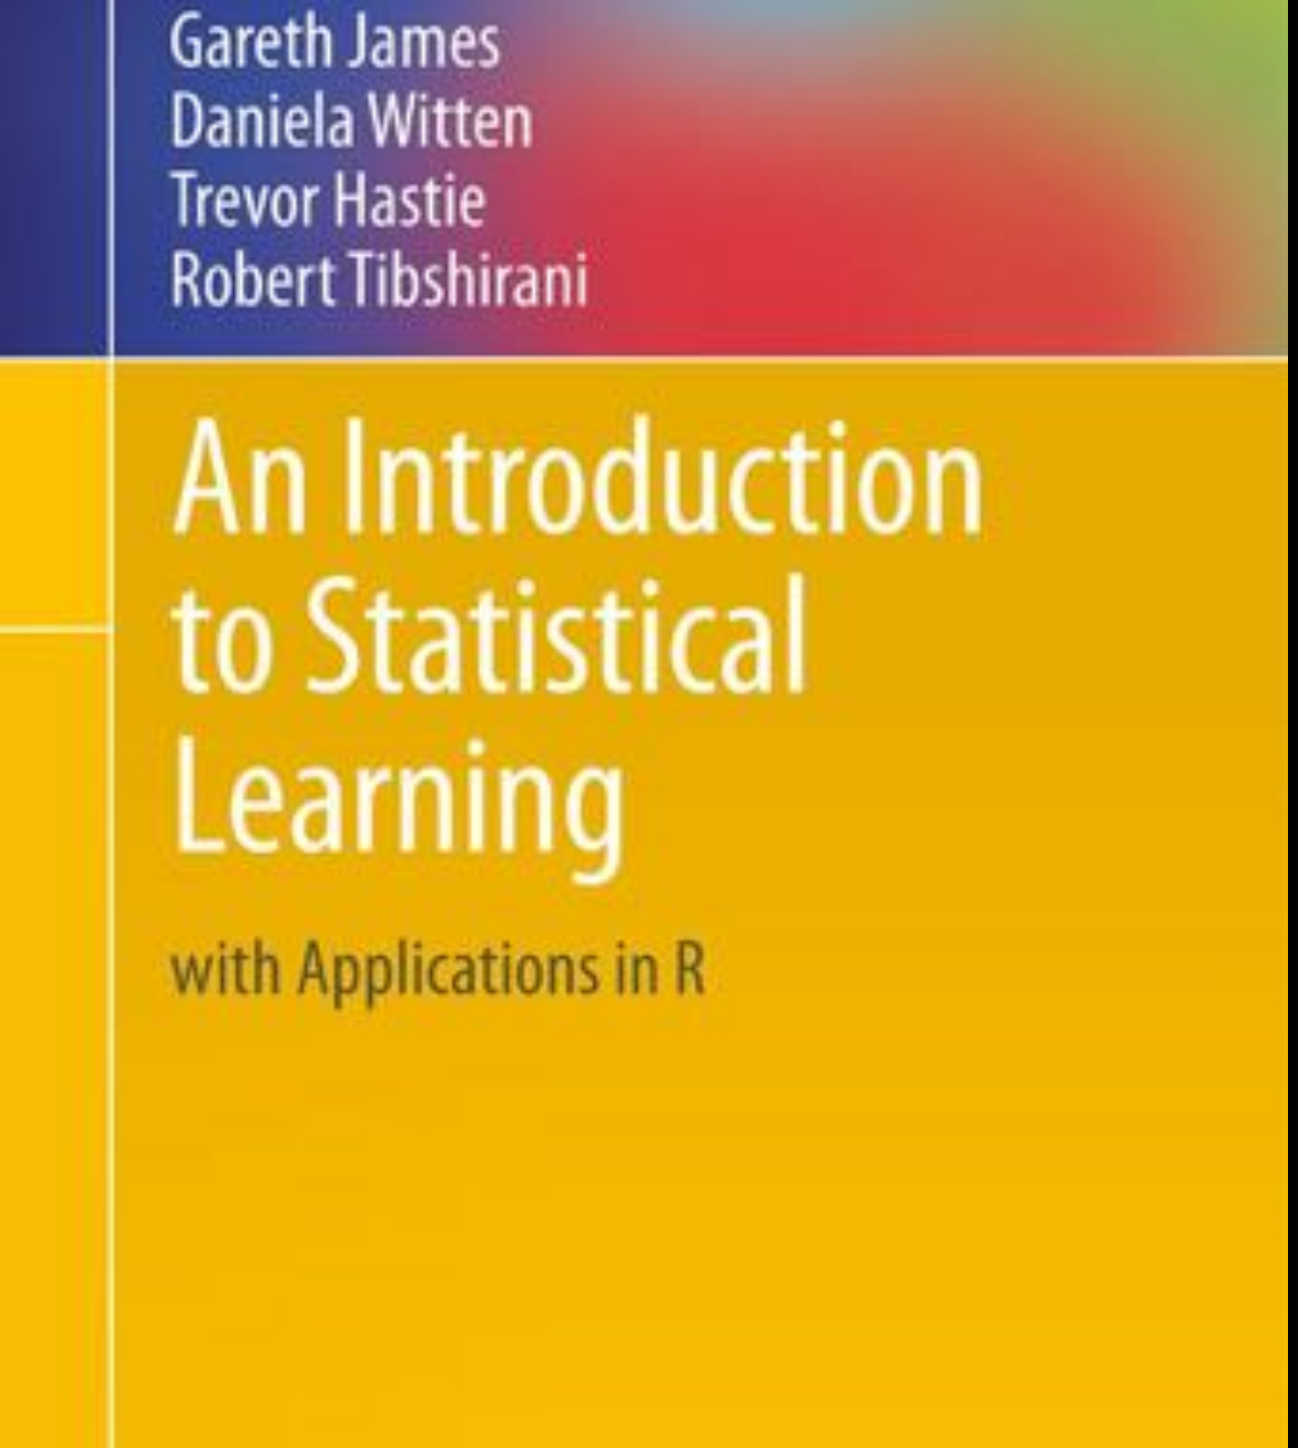

Resources - Statistical Learning

- One of the good books on statistical learning is ISLR -> An Introduction to Statistical Learning with Application in R
- You can find all the ISLR code in python at this github repo - https://github.com/JWarmenhoven/ISLR-python

Resources - Time Series

- Forecasting: Principle and Text
- Statistical forecasting: Notes on regression and time series analysis Case

Resources - Text Analytics

- Natural Language Processing with Python

Online Course

- Harvard Data Science Course - CS 109 Course (It is structured in similar way to the approach we shared)
- Data Science Specialisation - JHU Data Science (It is a good course, though the material is coded in R)
- Many more on Coursera & Udacity...


Speak to Us!


Thank you
@amitkaps | @bargava